

Détection d'événements à partir de capteurs sols – application au suivi de personnes fragiles

Soutenance de thèse

Ludovic Minvielle

Thèse CIFRE entre le Centre Borelli (ENS Paris-Saclay) et Tarkett

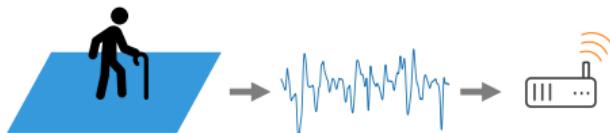
Mercredi 15 Juillet 2020



Introduction

General context

- ▶ Elderly population is growing (World Health Organization [6])
- ▶ Higher levels of frailty globally
- ▶ Increasing demand for reliable monitoring devices
- ▶ Tarkett, French company: 12,500 employees, 13 industrial sites, 1.3 millions m² of flooring per day
- ▶ *Floor in Motion*: a floor-based sensor for elderly care



Introduction

General context

- ▶ Elderly population is growing (World Health Organization [6])
- ▶ Higher levels of frailty globally
- ▶ Increasing demand for reliable monitoring devices
- ▶ Tarkett, French company: 12,500 employees, 13 industrial sites, 1.3 millions m² of flooring per day
- ▶ *Floor in Motion*: a floor-based sensor for elderly care

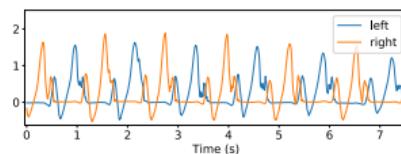


Tarkett's objective

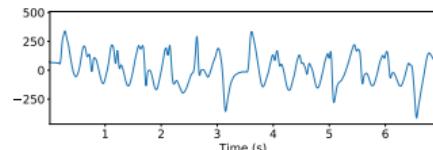
Providing tools for elderly monitoring
in nursing homes

Challenges

- ▶ Signal interpretability (external perturbations, artefacts, unidimensional for one area)
- ▶ Real world application (data scarcity, model for embedded systems)



Walk recorded with accelerometers



Walk recorded with Tarkett's floor sensor

Outline

1. Monitoring systems for fall detection
2. Fall detection using a floor sensor and machine learning
3. Transfer learning from experimental setup to operational data
4. Elderly activity recognition with convolutional neural networks
5. Conclusion

Monitoring systems for fall detection

Sensors

What makes a good monitoring system ?

- ▶ coverage and occlusion
- ▶ intrusiveness
- ▶ signal quality / information
- ▶ robustness
- ▶ ease of installation / use
- ▶ scalability

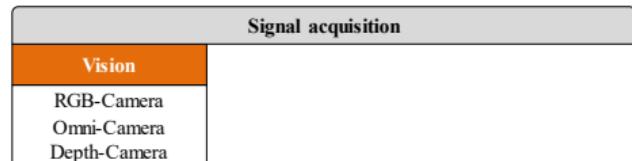
Criteria

Coverage/Occlusion
Intrusiveness
Signal quality / info
Robustness
Ease of instal. / use
Scalability

Sensors

What makes a good monitoring system ?

- ▶ coverage and occlusion
- ▶ intrusiveness
- ▶ signal quality / information
- ▶ robustness
- ▶ ease of installation / use
- ▶ scalability



Criteria	RGB cam	Depth cam
Coverage/Occlusion	★☆☆	★☆☆
Intrusiveness	★☆☆	★☆☆
Signal quality / info	★★★	★★★
Robustness	★☆☆	★★★
Ease of instal. / use	★☆☆	★☆☆
Scalability	★☆☆	★☆☆

Sensors

What makes a good monitoring system ?

- ▶ coverage and occlusion
- ▶ intrusiveness
- ▶ signal quality / information
- ▶ robustness
- ▶ ease of installation / use
- ▶ scalability

Signal acquisition	
Vision	Wearable
RGB-Camera	
Omni-Camera	Accelerometer
Depth-Camera	Gyroscope
	Barometric pressure

Criteria	RGB cam	Depth cam	Wearable
Coverage/Occlusion	★☆☆	★☆☆	★★★
Intrusiveness	★☆☆	★☆☆	★★☆
Signal quality / info	★★★	★★★	★★☆
Robustness	★☆☆	★★★	★★★
Ease of instal. / use	★☆☆	★☆☆	★★☆
Scalability	★☆☆	★☆☆	★★★

Sensors

What makes a good monitoring system ?

- ▶ coverage and occlusion
- ▶ intrusiveness
- ▶ signal quality / information
- ▶ robustness
- ▶ ease of installation / use
- ▶ scalability

Signal acquisition		
Vision	Wearable	Ambient
RGB-Camera	Accelerometer	Microphone
Omni-Camera	Gyroscope	Radar
Depth-Camera	Barometric pressure	Wi-Fi
		Vibration
		Floor

Criteria	RGB cam	Depth cam	Wearable	Acoustic	Radar / Wi-Fi	Vibration	Floor
Coverage/Occlusion	★☆☆	★☆☆	★★★	★☆☆	★☆☆	★★★	★★★
Intrusiveness	★☆☆	★☆☆	★☆☆	★☆☆	★☆☆	★★★	★★★
Signal quality / info	★★★	★★★	★☆☆	★☆☆	★☆☆	★☆☆	★☆☆
Robustness	★☆☆	★★★	★★★	★☆☆	★☆☆	★☆☆	★☆☆
Ease of instal. / use	★☆☆	★☆☆	★☆☆	★☆☆	★☆☆	★★★	★☆☆
Scalability	★☆☆	★☆☆	★★★	★☆☆	★☆☆	★☆☆	★★★

Information extraction

How to process the inputs ?

- ▶ All systems use feature extraction
- ▶ The “level” of feature engineering depends on the complexity / dimensionality of the input signal

How to deal with feature signals ?

- ▶ Use simple thresholds
- ▶ Use them as feature vectors for classification models (anomaly detection, classical supervised models)

Signal acquisition		
Vision	Wearable	Ambient
RGB-Camera	Accelerometer	Microphone
Omni-Camera	Gyroscope	Radar
Depth-Camera	Barometric pressure	Wi-Fi
		Vibrational
		Floor



Feature extraction		
Vision	Wearable	Ambient
Position	Position	Statistical measures
Motion	Velocity	Fourier transform
Shape	Angle	Wavelet transform
		Cepstrum features



Decision rule	
Threshold	Machine learning
	kNN SVM HMM Decision Tree

Fall detection using a floor sensor and machine learning

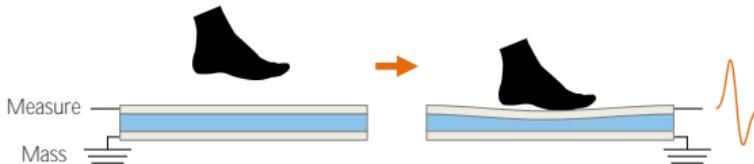
Tarkett sensor

- ▶ Piezoelectric principle:

$$d = \frac{Q}{F},$$

(simple version) with d the *piezoelectric constant*.

When stressed or squeezed, the material emits charges.



Tarkett sensor

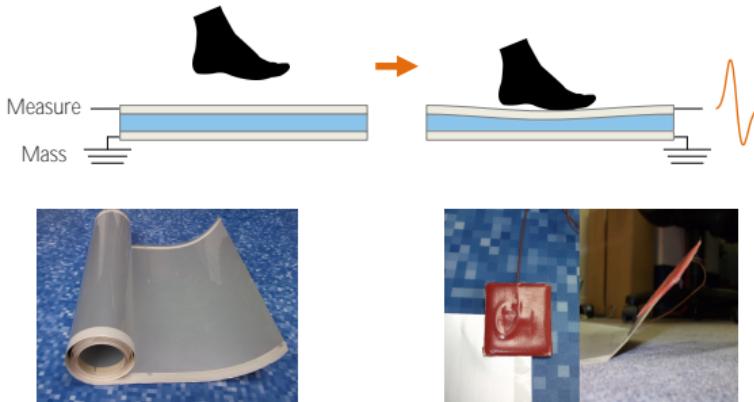
- ▶ Piezoelectric principle:

$$d = \frac{Q}{F},$$

(simple version) with d the *piezoelectric constant*.

When stressed or squeezed, the material emits charges.

- ▶ How does this look like ?
0.3 mm thick and 60 cm wide roll
with customizable length



Tarkett sensor

- ▶ Piezoelectric principle:

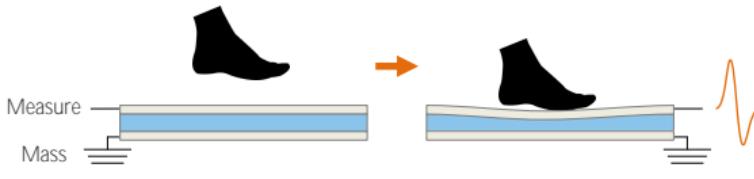
$$d = \frac{Q}{F},$$

(simple version) with d the *piezoelectric constant*.

When stressed or squeezed, the material emits charges.

- ▶ How does this look like ?
0.3 mm thick and 60 cm wide roll
with customizable length
- ▶ How is it installed ?

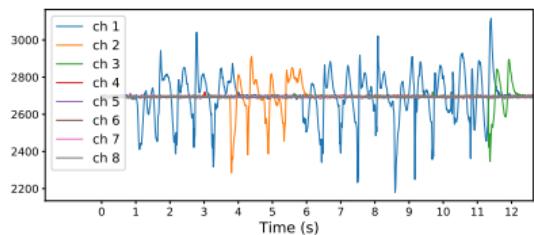
- ▶ Under the flooring
- ▶ Several connected bands for each area, hence one area corresponds to one input



Data

Preprocessing

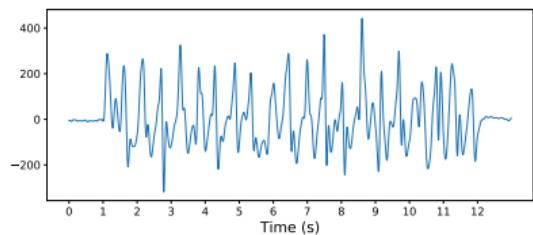
- ▶ linear detrending
- ▶ low-pass filtering
- ▶ zeroing low energy channels
- ▶ sum over all channels



Data

Preprocessing

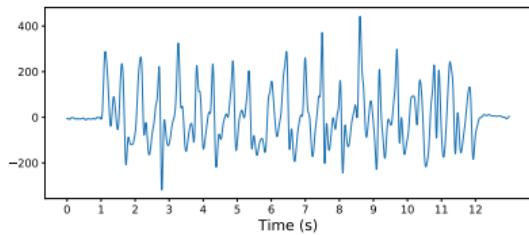
- ▶ linear detrending
- ▶ low-pass filtering
- ▶ zeroing low energy channels
- ▶ sum over all channels



Data

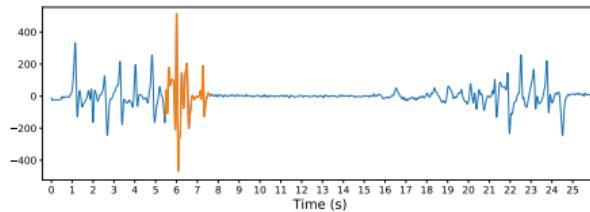
Preprocessing

- ▶ linear detrending
- ▶ low-pass filtering
- ▶ zeroing low energy channels
- ▶ sum over all channels



Experimental dataset

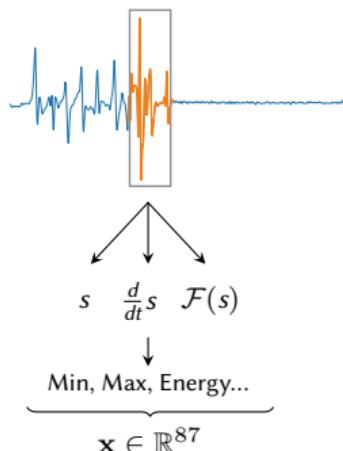
- ▶ 28 volunteers aged 25 to 45
- ▶ 742 signals collected in **controlled environment**
- ▶ 55% fall, 45% non-fall
- ▶ varied fall events (forward, backward...) and activities of daily living (walking, sitting...)



Model

Time series as *feature vector*. At every timestamp:

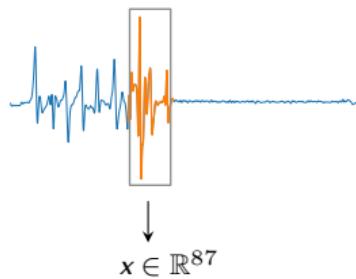
1. Window over the signal: 2.5 s
2. Compute feature vector: 29 statistical measures
(Min, Max, Shannon energy, Percentile,...) over
three representations of the signal



Model

Time series as *feature vector*. At every timestamp:

1. Window over the signal: 2.5 s
2. Compute feature vector: 29 statistical measures (Min, Max, Shannon energy, Percentile,...) over three representations of the signal



3. Classification model: Random Forest (Breiman [1]), based on **decision trees**

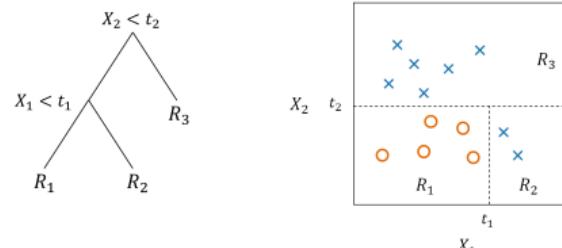
Decision tree

Feature space $\mathcal{X} = \mathbb{R}^Q$. Division of \mathcal{X} into non-overlapping regions R_1, \dots, R_j . Algorithm CART: recursive binary splits [2] that solve:

$$\arg \max_{X_q, \tau} \text{IG} , \quad (\text{information gain})$$

$$\text{with } \text{IG}(X_q, \tau) = I(n) - \frac{N_l}{N_n} I(l) - \frac{N_r}{N_n} I(r) ,$$

$$\text{and } I(n) = \text{Gini}(n) = \sum_k p_{nk}(1 - p_{nk}) .$$



$$\text{Prediction function: } f(x) = \sum_{j=1}^J c_j \mathbb{1}(x \in R_j)$$

Model

Random forest

Decision trees d_1, \dots, d_{N_T} grown with two rules:

- ▶ Each tree is trained with a *bootstrap* of the training set
- ▶ At each split, access to a random subset of pool of features

Each tree is a “vote” for a class. The estimated probability of belonging to class k is then:

$$f_k(x) = \frac{1}{N_T} \sum_{i=1}^{N_T} \mathbb{1}(d_i(x) = k)$$

Model

Random forest

Decision trees d_1, \dots, d_{N_T} grown with two rules:

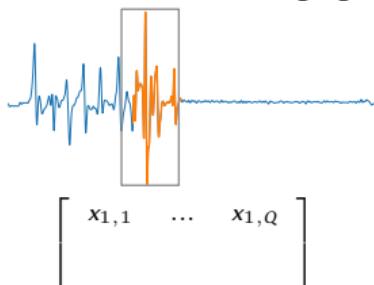
- ▶ Each tree is trained with a *bootstrap* of the training set
- ▶ At each split, access to a random subset of pool of features

Each tree is a “vote” for a class. The estimated probability of belonging to class k is then:

$$f_k(x) = \frac{1}{N_T} \sum_{i=1}^{N_T} \mathbb{1}(d_i(x) = k)$$

Data augmentation

Select r windows in training signals



- ▶ *Fall* signals: encompass the fall
- ▶ *Non-fall* signals: random location

Model

Random forest

Decision trees d_1, \dots, d_{N_T} grown with two rules:

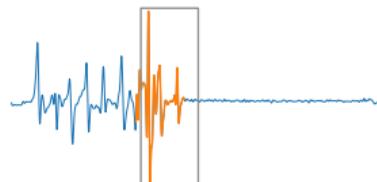
- ▶ Each tree is trained with a *bootstrap* of the training set
- ▶ At each split, access to a random subset of pool of features

Each tree is a “vote” for a class. The estimated probability of belonging to class k is then:

$$f_k(x) = \frac{1}{N_T} \sum_{i=1}^{N_T} \mathbb{1}(d_i(x) = k)$$

Data augmentation

Select r windows in training signals



$$\begin{bmatrix} x_{1,1} & \dots & x_{1,Q} \\ x_{2,1} & \dots & x_{2,Q} \end{bmatrix}$$

- ▶ *Fall* signals: encompass the fall
- ▶ *Non-fall* signals: random location

Model

Random forest

Decision trees d_1, \dots, d_{N_T} grown with two rules:

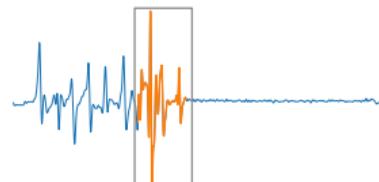
- ▶ Each tree is trained with a *bootstrap* of the training set
- ▶ At each split, access to a random subset of pool of features

Each tree is a “vote” for a class. The estimated probability of belonging to class k is then:

$$f_k(x) = \frac{1}{N_T} \sum_{i=1}^{N_T} \mathbb{1}(d_i(x) = k)$$

Data augmentation

Select r windows in training signals



$$\begin{bmatrix} x_{1,1} & \dots & x_{1,Q} \\ x_{2,1} & \dots & x_{2,Q} \\ x_{3,1} & \dots & x_{3,Q} \end{bmatrix}$$

- ▶ *Fall* signals: encompass the fall
- ▶ *Non-fall* signals: random location

Model

Random forest

Decision trees d_1, \dots, d_{N_T} grown with two rules:

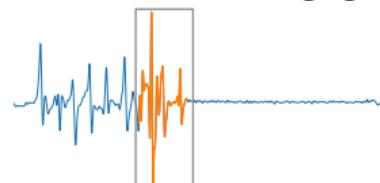
- ▶ Each tree is trained with a *bootstrap* of the training set
- ▶ At each split, access to a random subset of pool of features

Each tree is a “vote” for a class. The estimated probability of belonging to class k is then:

$$f_k(x) = \frac{1}{N_T} \sum_{i=1}^{N_T} \mathbb{1}(d_i(x) = k)$$

Data augmentation

Select r windows in training signals

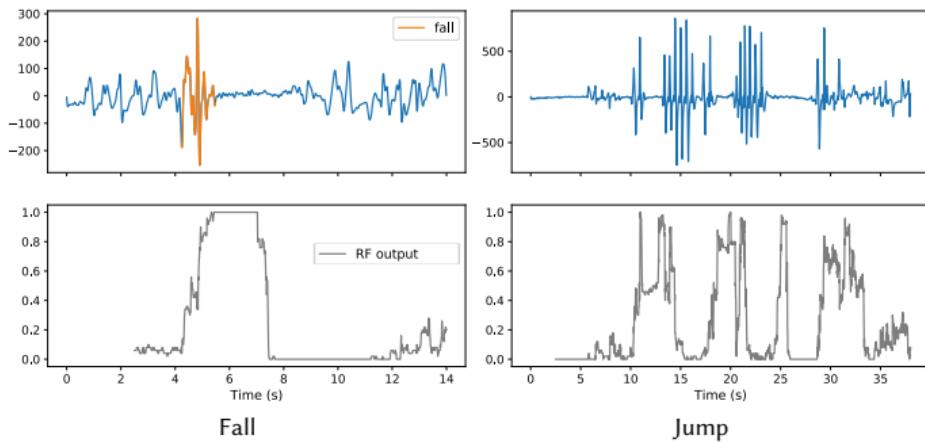


$$\begin{bmatrix} x_{1,1} & \dots & x_{1,Q} \\ \vdots & & \vdots \\ x_{r,1} & \dots & x_{r,Q} \end{bmatrix}$$

- ▶ *Fall* signals: encompass the fall
- ▶ *Non-fall* signals: random location

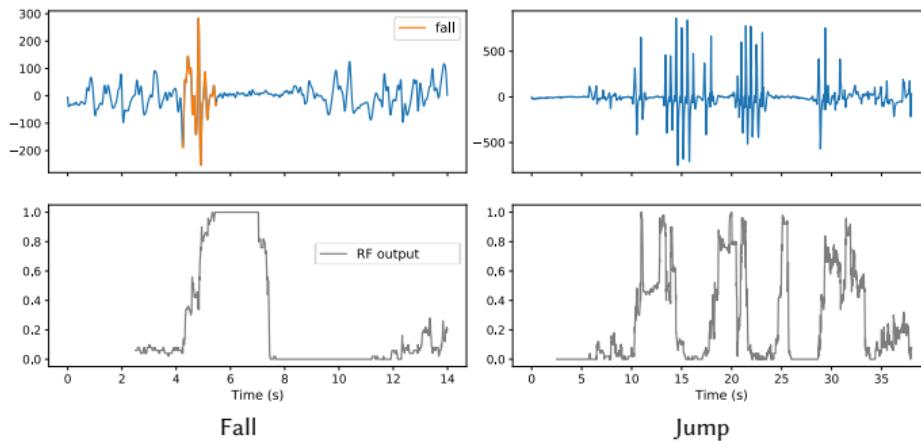
Model

Time aggregation



Model

Time aggregation



Time aggregation

$N_f(t)$: number of trees voting for *fall*

Use a buffer $B_s \in \mathbb{N}$ and a threshold $T_h \in [0, 1]$

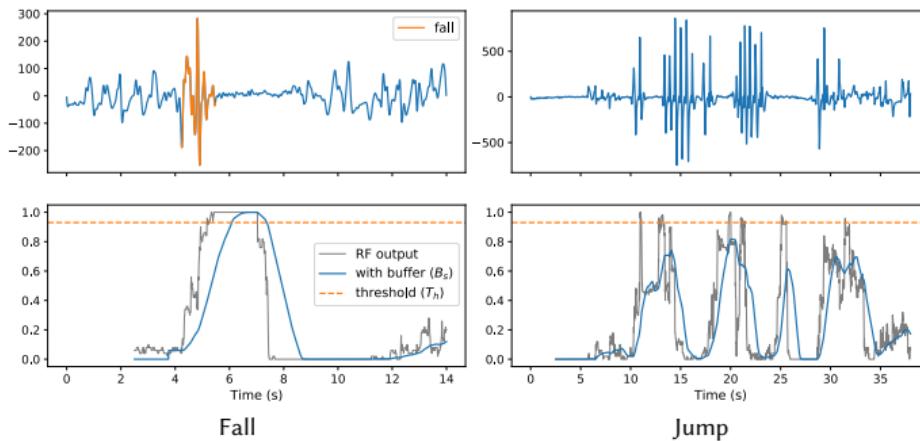
$$g(t) = \frac{\sum_{u=t-B_s+1}^t N_f(u)}{B_s \times N_T}$$

New binary classification function:

$$d(t) = \begin{cases} 1, & \text{if } g(t) > T_h \\ 0, & \text{otherwise} \end{cases}$$

Model

Time aggregation



Time aggregation

$N_f(t)$: number of trees voting for *fall*

Use a buffer $B_s \in \mathbb{N}$ and a threshold $T_h \in [0, 1]$

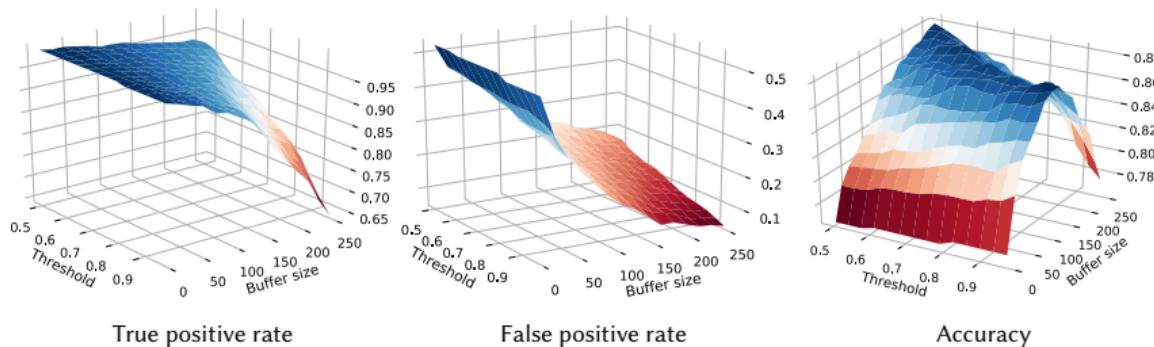
$$g(t) = \frac{\sum_{u=t-B_s+1}^t N_f(u)}{B_s \times N_T}$$

New binary classification function:

$$d(t) = \begin{cases} 1, & \text{if } g(t) > T_h \\ 0, & \text{otherwise} \end{cases}$$

Model

Parameters evaluation



True positive rate

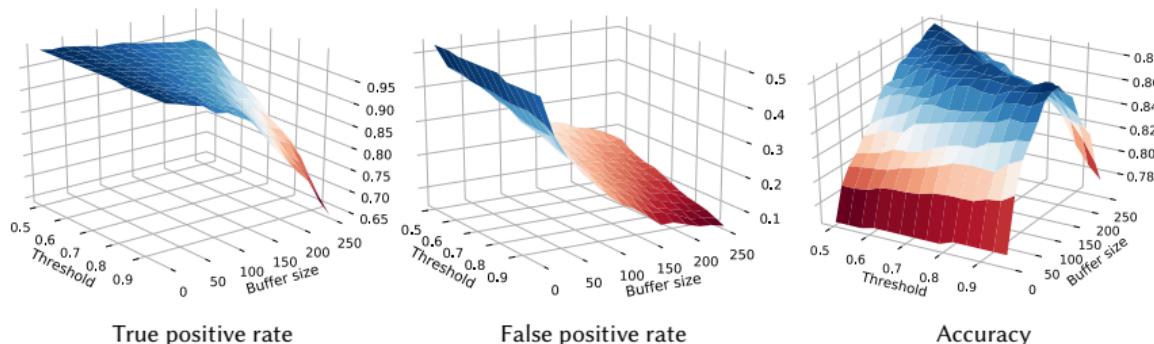
False positive rate

Accuracy

- ▶ Buffer/threshold trade-off useful for maintaining low FPR while improving TPR

Model

Parameters evaluation



- ▶ Buffer/threshold trade-off useful for maintaining low FPR while improving TPR

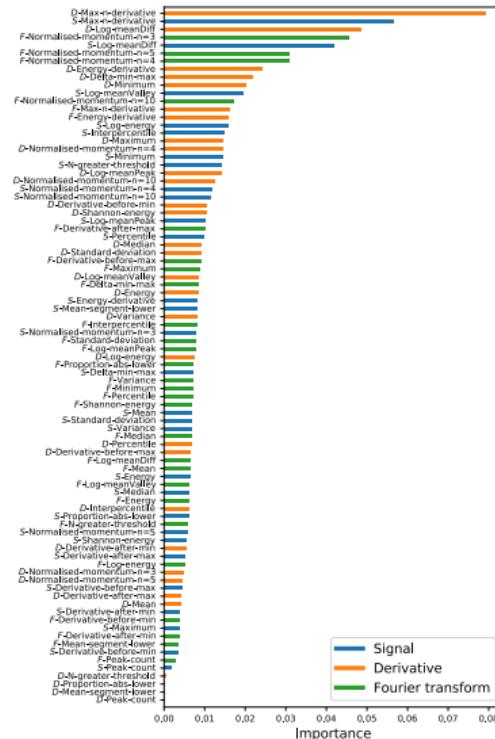
Decision rule is ready. Is it implementable in the local processing unit ?

Feature selection

Feature importance

$$\text{Tree: } I(X_q) = \sum_{\text{nodes } t} p(t) \Delta i(t) \mathbb{1}(v(t) = X_q)$$

$$\text{Random forest: } I(X_q) = \frac{1}{N_T} \sum_{n=1}^{N_T} I(T_n, X_q)$$



Feature selection

Feature importance

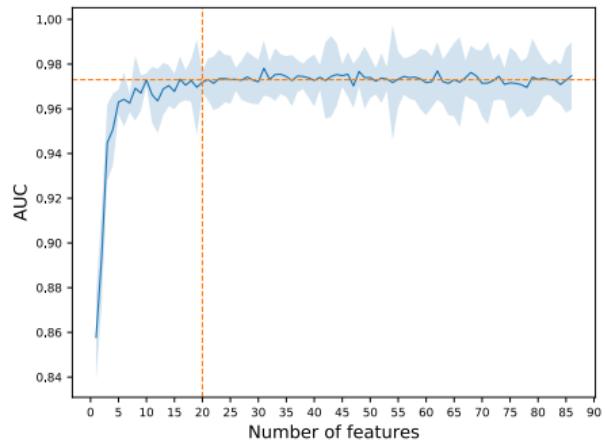
$$\text{Tree: } I(X_q) = \sum_{\text{nodes } t} p(t) \Delta i(t) \mathbb{1}(v(t) = X_q)$$

$$\text{Random forest: } I(X_q) = \frac{1}{N_T} \sum_{n=1}^{N_T} I(T_n, X_q)$$

Recursive feature elimination

Initial pool of Q features X_1, \dots, X_Q .

1. Train several times and record variable importances
2. Average of importances over trainings.
 $X_{q^*} = \arg \min_{X_i} I(X_i)$
3. Remove X_{q^*} from the pool of features and back to step 1



Results

Set up:

- ▶ Fixed params: r set to 5 and Q set to 20
- ▶ Varying T_h (0.5 to 1) and B_s (5 to 250)
- ▶ Record best Accuracy and show corresponding TPR, FPR

Model	Accuracy	TPR	FPR
LR	86.8 ± 1.5	90.5 ± 2.4	17.7 ± 4.9
LDA	85.5 ± 1.2	91.0 ± 2.1	21.7 ± 3.7
k-NN	87.0 ± 1.9	89.2 ± 1.4	16.0 ± 4.7
SVM	87.6 ± 3.2	90.0 ± 4.5	15.5 ± 6.8
MLP	88.2 ± 1.5	92.4 ± 1.2	17.3 ± 4.1
RF	88.2 ± 1.5	91.7 ± 3.5	16.2 ± 6.2

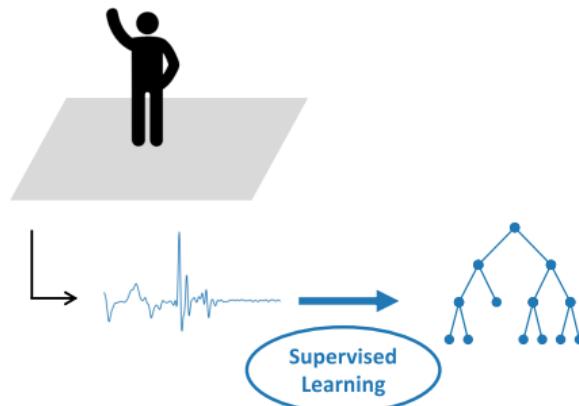
Comments:

- ▶ Parametric methods perform worse than non-parametric
- ▶ RF is slightly better than others

Transfer learning from experimental setup to operational data

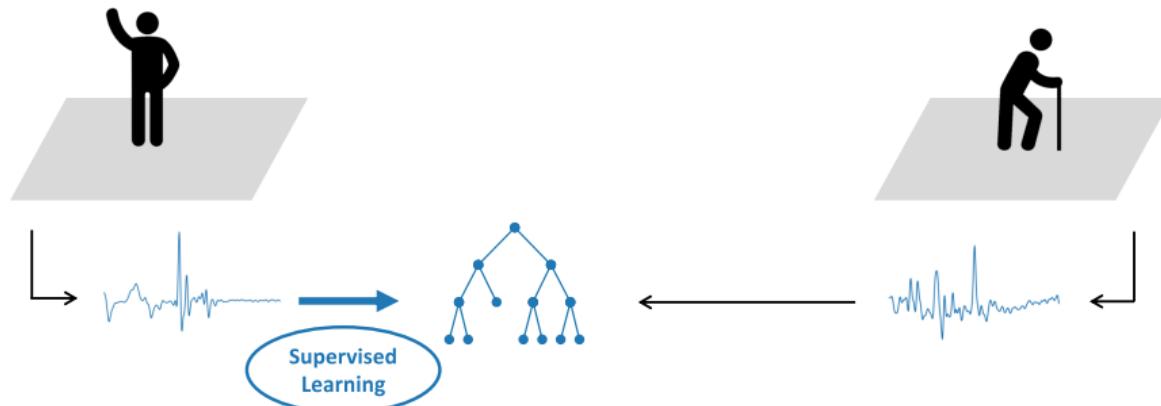
Context

Experimental vs. operational



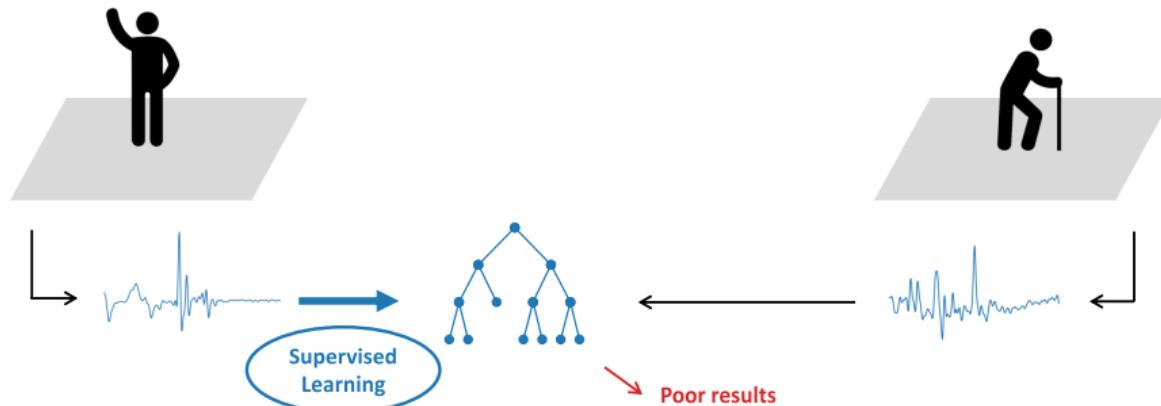
Context

Experimental vs. operational



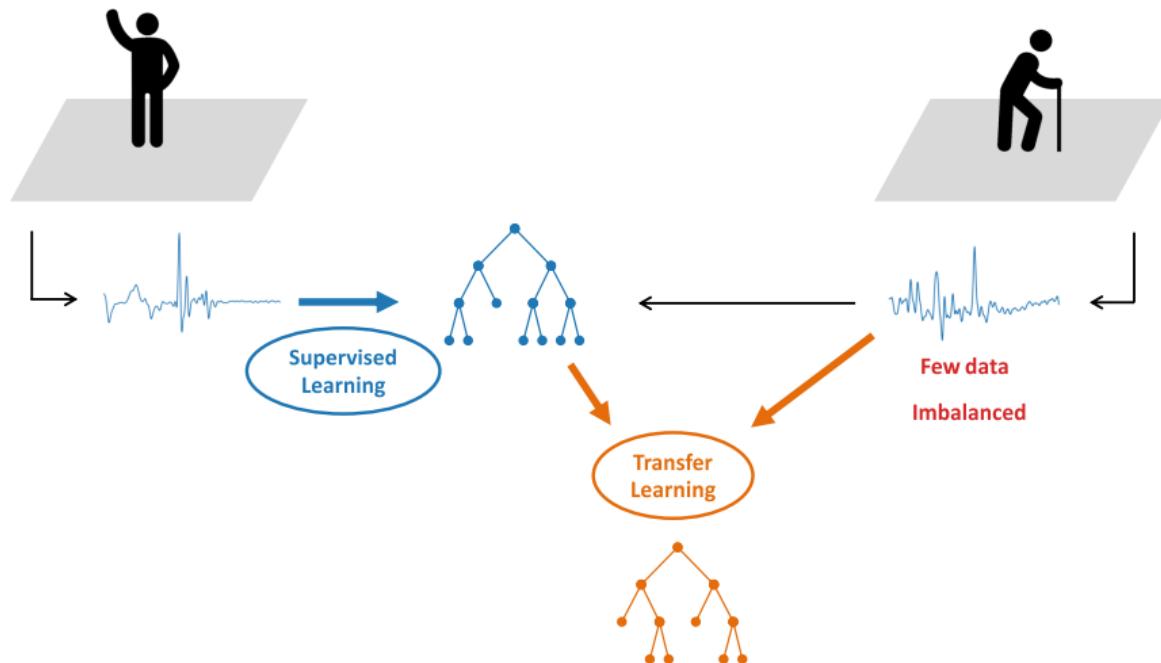
Context

Experimental vs. operational



Context

Experimental vs. operational



Context

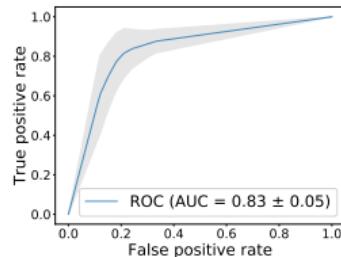
Transfer learning

Transfer learning

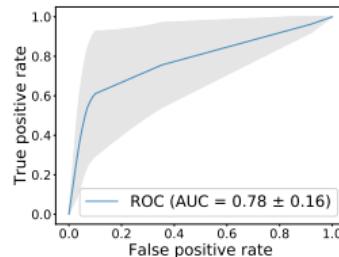
- ▶ Source domain: $\mathcal{D}_S = \{\mathcal{X}_S, P(X_S)\}$
- ▶ Target domain: $\mathcal{D}_T = \{\mathcal{X}_T, P(X_T)\}$
- ▶ Source task: $\mathcal{T}_S = \{\mathcal{Y}_S, f^S\}$
- ▶ Target task: $\mathcal{T}_T = \{\mathcal{Y}_T, f^T\}$

Our case

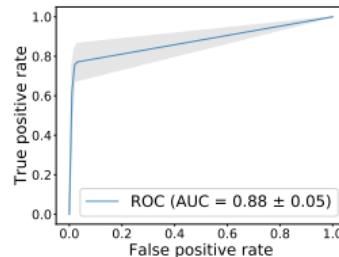
- ▶ $\mathcal{X}_S = \mathcal{X}_T$
- ▶ $P(X_S) \neq P(X_T)$
- ▶ $\mathcal{Y}_S = \mathcal{Y}_T$
- ▶ $f^S \neq f^T$



Source tested on source



Source tested on target



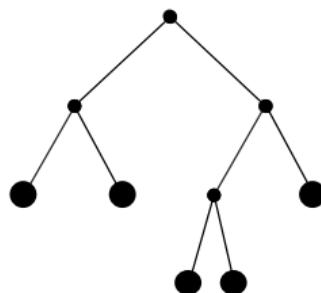
Target tested on target

- ▶ **Goal:** Use knowledge from source and target domains to improve the final task while avoiding *negative transfer*
- ▶ Model-based transfer: we have access to the Source model and target data

Model-based transfer

Segev et al. [5]

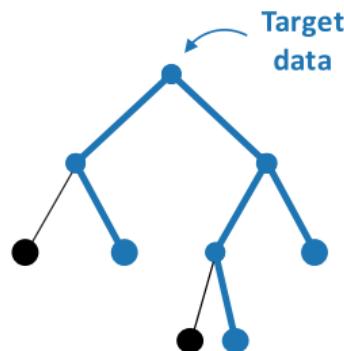
Structure Expansion / Reduction (SER)



Model-based transfer

Segev et al. [5]

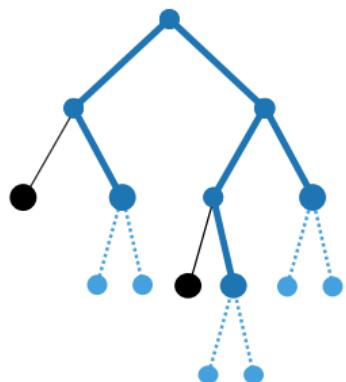
Structure Expansion / Reduction (SER)



Model-based transfer

Segev et al. [5]

Structure Expansion / Reduction (SER)

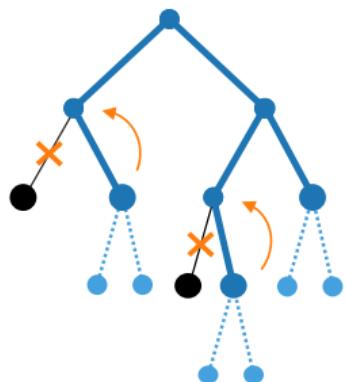


1. Expansion

Model-based transfer

Segev et al. [5]

Structure Expansion / Reduction (SER)

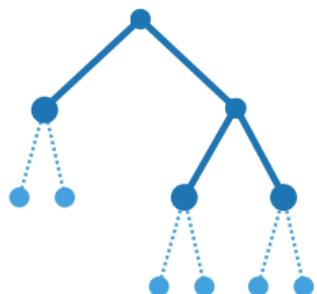


1. Expansion
2. Reduction

Model-based transfer

Segev et al. [5]

Structure Expansion / Reduction (SER)

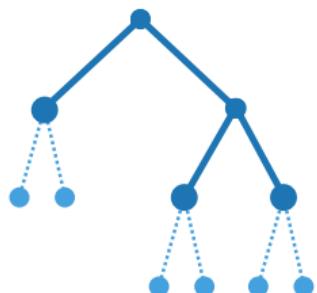


1. Expansion
2. Reduction

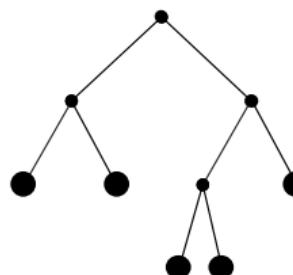
Model-based transfer

Segev et al. [5]

Structure Expansion / Reduction (SER)



Structure Transfer (STRUT)

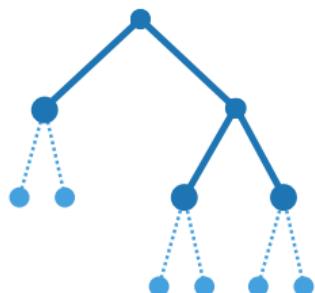


1. Expansion
2. Reduction

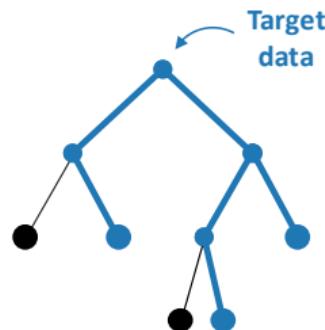
Model-based transfer

Segev et al. [5]

Structure Expansion / Reduction (SER)



Structure Transfer (STRUT)

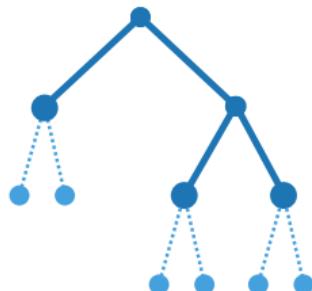


1. Expansion
2. Reduction

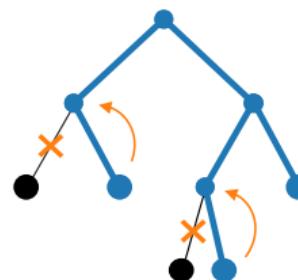
Model-based transfer

Segev et al. [5]

Structure Expansion / Reduction (SER)



Structure Transfer (STRUT)



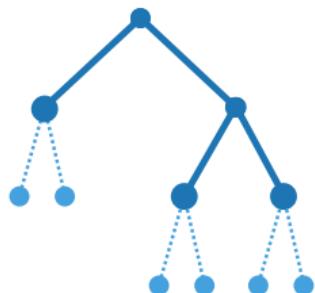
1. Pruning

1. Expansion
2. Reduction

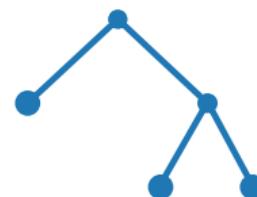
Model-based transfer

Segev et al. [5]

Structure Expansion / Reduction (SER)



Structure Transfer (STRUT)



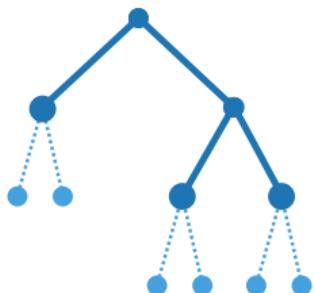
1. Pruning

1. Expansion
2. Reduction

Model-based transfer

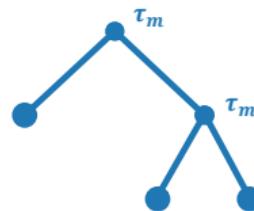
Segev et al. [5]

Structure Expansion / Reduction (SER)



1. Expansion
2. Reduction

Structure Transfer (STRUT)

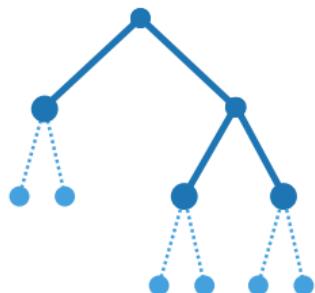


1. Pruning
2. Threshold update

Model-based transfer

Segev et al. [5]

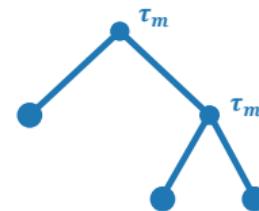
Structure Expansion / Reduction (SER)



1. Expansion
2. Reduction

Partition refinement or simplification

Structure Transfer (STRUT)



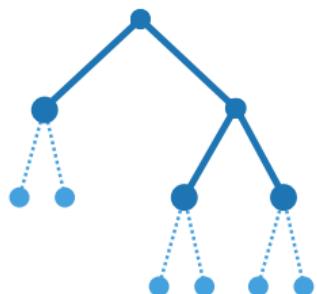
1. Pruning
2. Threshold update

Drifts

Model-based transfer

Segev et al. [5]

Structure Expansion / Reduction (SER)

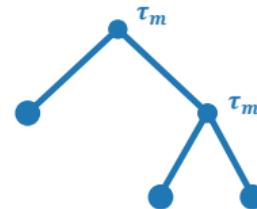


1. Expansion
2. Reduction

Partition refinement or simplification

How about imbalanced data ?

Structure Transfer (STRUT)



1. Pruning
2. Threshold update

Drifts

Leaf loss risk

Homogeneous class imbalance

$$P^T(x|y) = P^S(x|y)$$

$$P^T(y|x) = \lambda_y \frac{P^S(y|x)}{\int \lambda_y P^S(y|x) dy}$$

$$\text{with } \lambda_y = \frac{P^T(y)}{P^S(y)}$$

Leaf loss risk

Homogeneous class imbalance

$$P^T(x|y) = P^S(x|y)$$

$$P^T(y|x) = \lambda_y \frac{P^S(y|x)}{\int \lambda_y P^S(y|x) dy}$$

$$\text{with } \lambda_y = \frac{P^T(y)}{P^S(y)}$$

Leaf loss risk

Significant leaf: Leaf l that conserves the minority class k_{min} after Target update:

$$\forall k \neq k_{min}, \quad P^T(y = k_{min} | x \in l) > P^T(y = k | x \in l)$$

Leaf loss risk:

$$R_l(l) = P^T(x \notin l | y = k_{min})^{n_{k_{min}}}$$

Leaf loss risk

Homogeneous class imbalance

$$P^T(x|y) = P^S(x|y)$$

$$P^T(y|x) = \lambda_y \frac{P^S(y|x)}{\int \lambda_y P^S(y|x) dy}$$

with $\lambda_y = \frac{P^T(y)}{P^S(y)}$

Leaf loss risk

Significant leaf: Leaf l that conserves the minority class k_{min} after Target update:

$$\forall k \neq k_{min}, \quad P^T(y = k_{min} | x \in l) > P^T(y = k | x \in l)$$

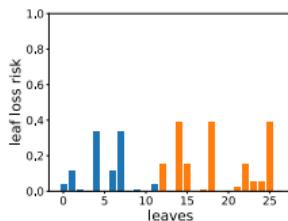
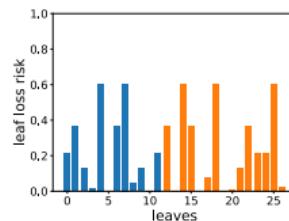
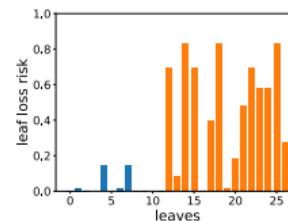
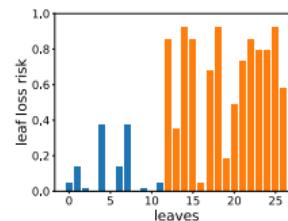
Leaf loss risk:

$$R_L(l) = P^T(x \notin l | y = k_{min})^{n_{k_{min}}}$$

Leaf loss risk under homogeneous class imbalance

$$\forall k \neq k_{min}, \quad \lambda_{k_{min}} P^S(y = k_{min} | x \in l) > \lambda_k P^S(y = k | x \in l)$$

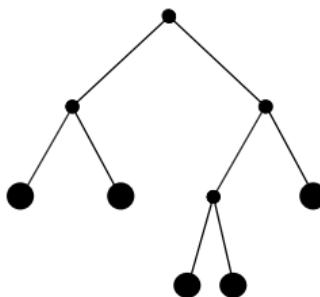
$$R_L(l) = P^S(x \notin l | y = k_{min})^{n_{k_{min}}}$$

(a) Balanced data, $n = 200$ (b) Balanced data, $n = 100$ (c) Imbalanced data (10%) $n = 200$ (d) Imbalanced data (10%) $n = 100$

SER for class imbalance

SER_R, SER_{LL}

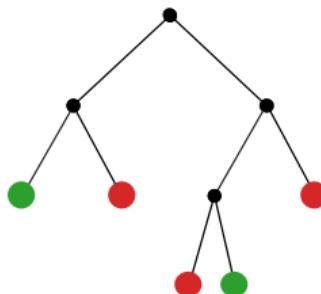
Structure Expansion and controlled Reduction



SER for class imbalance

SER_R, SER_{LL}

Structure Expansion and controlled Reduction

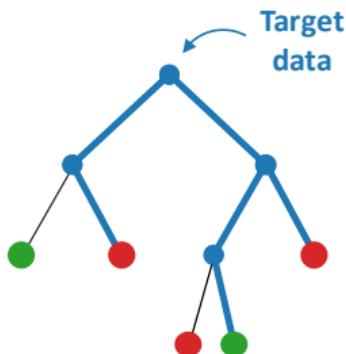


Minority class

SER for class imbalance

SER_R, SER_{LL}

Structure Expansion and controlled Reduction

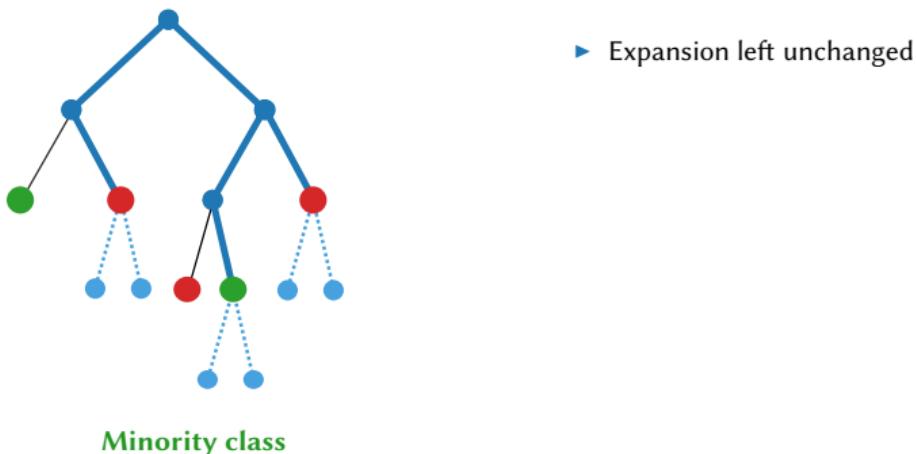


Minority class

SER for class imbalance

SER_R, SER_{LL}

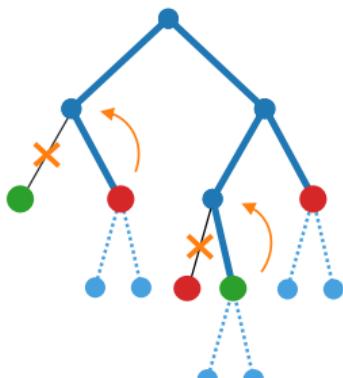
Structure Expansion and controlled Reduction



SER for class imbalance

SER_R, SER_{LL}

Structure Expansion and controlled Reduction



Minority class

- ▶ Expansion left unchanged
- ▶ Reduction constrained

SER_R

If node is of minority class, then no pruning

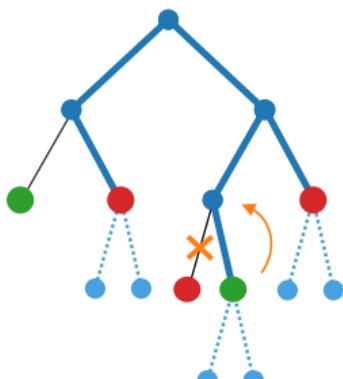
SER_{LL}

If node is of minority class **and** still significant considering Target **and** $R_L > 0.5$, then no pruning

SER for class imbalance

SER_R, SER_{LL}

Structure Expansion and controlled Reduction



- ▶ Expansion left unchanged
- ▶ Reduction constrained

SER_R

If node is of minority class, then no pruning

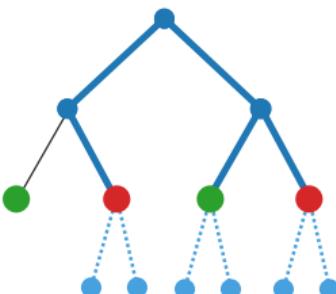
SER_{LL}

If node is of minority class **and** still significant considering Target **and** $R_L > 0.5$, then no pruning

SER for class imbalance

SER_R, SER_{LL}

Structure Expansion and controlled Reduction



Minority class

- ▶ Expansion left unchanged
- ▶ Reduction constrained

SER_R

If node is of minority class, then no pruning

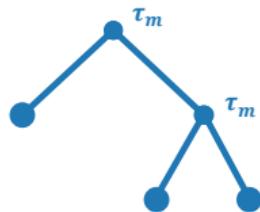
SER_{LL}

If node is of minority class **and** still significant considering Target **and** $R_L > 0.5$, then no pruning

STRUT for class imbalance

STRUT optimization

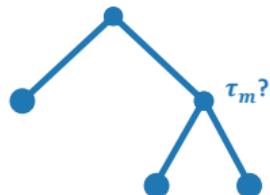
STRUT: How are updated the new thresholds ?



STRUT for class imbalance

STRUT optimization

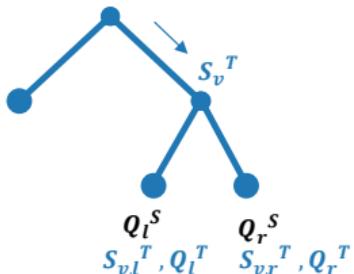
STRUT: How are updated the new thresholds ?



STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?



Q_l^S, Q_r^S : class proportions of source data in children w.r.t. the *original* split

$S_{v,l}^T, S_{v,r}^T$: subsets of S_v^T that fall in the children nodes of v

$Q_l^T(\tau), Q_r^T(\tau)$: class proportions of target data in children w.r.t. the *new* split

Divergence Gain: similarity between the original label distributions and the new ones

$$DG(\tau) = 1 - \frac{|S_{v,l}^T|}{|S_v^T|} JSD(Q_l^S, Q_l^T) - \frac{|S_{v,r}^T|}{|S_v^T|} JSD(Q_r^S, Q_r^T)$$

Jensen-Shannon divergence:

$$JSD(P, Q) = \frac{1}{2} (D_{KL}(P||M) + D_{KL}(Q||M))$$

$$M = \frac{1}{2} (P + Q)$$

Kullback-Leibler divergence:

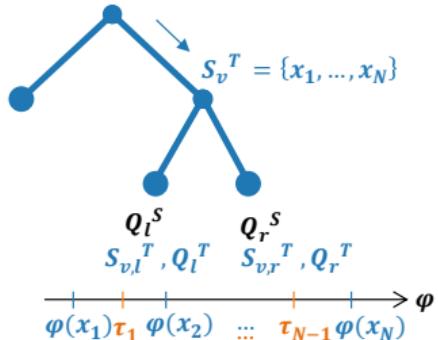
$$D_{KL}(P||Q) = \sum_k P(k) \ln \left(\frac{P(k)}{Q(k)} \right)$$

STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?

Goal: Maximize DG while being in a local maximum of Information Gain (IG) (here IG = Gini gain)



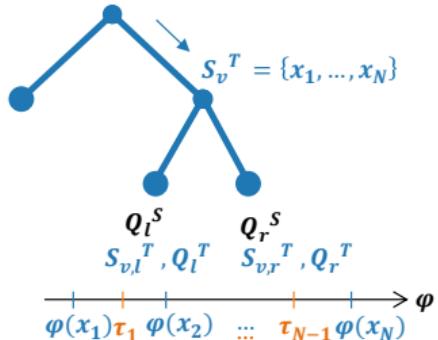
$$\tau_m = \arg \max_{\tau \in T_v} (DG(\tau, Q_l^T(\tau), Q_r^T(\tau)))$$

$$\text{s.t. } IG(\tau_{m-1}) < IG(\tau_m) \text{ and } IG(\tau_m) > IG(\tau_{m+1})$$

STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?



Goal: Maximize DG while being in a local maximum of Information Gain (IG) (here IG = Gini gain)

$$\tau_m = \arg \max_{\tau \in T_v} (DG(\tau, Q_l^T(\tau), Q_r^T(\tau)))$$

$$\text{s.t. } IG(\tau_{m-1}) < IG(\tau_m) \text{ and } IG(\tau_m) > IG(\tau_{m+1})$$

- ▶ Q^S have less meaning when going deeper
- ▶ Do we really want to keep Q^S and Q^T close ?

STRUT for class imbalance

STRUT_{IG}, STRUT_{HI}

STRUT_{IG}

STRUT without DG: maximization of IG

STRUT for class imbalance

STRUT_{IG}, STRUT_{HI}

STRUT_{IG}

STRUT without DG: maximization of IG

STRUT_{HI}

Framework of homogeneous class imbalance, use

$$p^T(y/x) = \lambda_y \frac{p^S(y/x)}{\int \lambda_y p^S(y/x) dy}$$

to change the source class proportions in DG:

$$Q_l^{S'} = \lambda_k \frac{Q_l^S}{\sum_k \lambda_k Q_l^S} \quad Q_r^{S'} = \lambda_k \frac{Q_r^S}{\sum_k \lambda_k Q_r^S}$$

STRUT for class imbalance

STRUT_{IG}, STRUT_{HI}

STRUT_{IG}

STRUT without DG: maximization of IG

STRUT_{HI}

Framework of homogeneous class imbalance, use

$$p^T(y/x) = \lambda_y \frac{p^S(y/x)}{\int \lambda_y p^S(y/x) dy}$$

to change the source class proportions in DG:

$$Q_l^{S'} = \lambda_k \frac{Q_l^S}{\sum_k \lambda_k Q_l^S} \quad Q_r^{S'} = \lambda_k \frac{Q_r^S}{\sum_k \lambda_k Q_r^S}$$

STRUT_{HI} can be seen as a generalization of STRUT

Results

Synthetic data

Gaussian generator

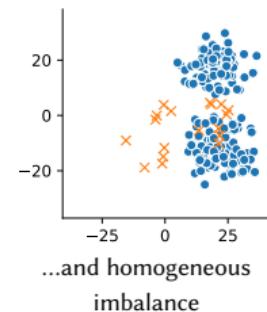
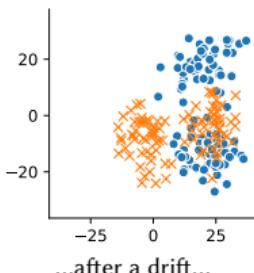
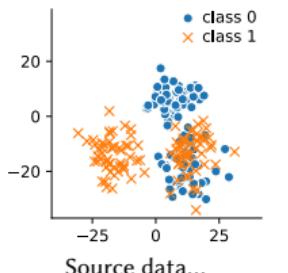
Source dataset: combination of several multivariate Gaussian clusters. We use $N_{\text{source}} = 200$ and $N_{\text{clust}} = 10$. Initial parameters are randomly drawn from Uniform distribution: $\mu_i \in [-70, 70]$ and $\sigma_i \in [5, 15]$

Transformations

Basic transformations on Source clusters are applied to get a Target dataset:

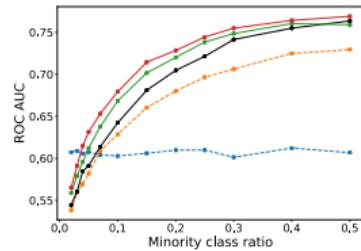
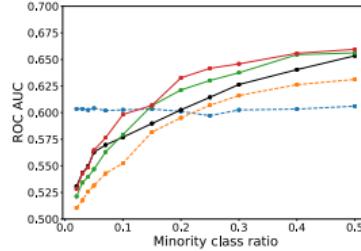
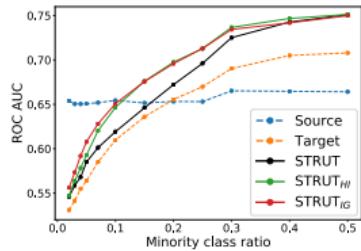
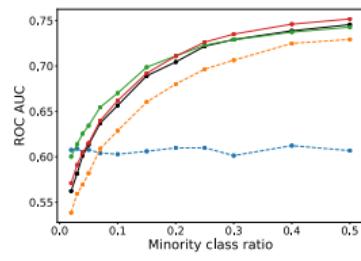
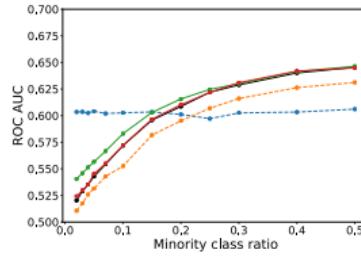
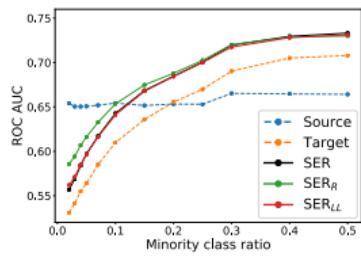
- ▶ Drifts (μ_i)
- ▶ Stretch / Squeeze (σ_i)
- ▶ Adding / Remove clusters

Transformations are combined with imbalance ratio (from 2% to 50%), and $N_{\text{target}} = 1000$



Results

Synthetic data



(a) Drift

(b) Stretch / Squeeze

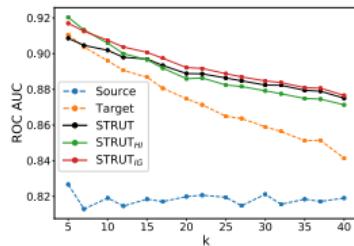
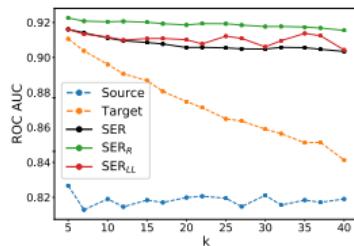
(c) Add / Remove

- ▶ Strong imbalance leads to negative transfer
- ▶ SER: SER_{LL} close to SER, SER_R better than SER (esp. with high imbalance)
- ▶ STRUT: both variants beat the original STRUT

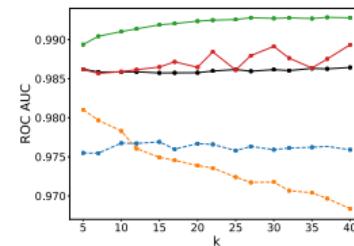
Results

Fall data

Operational data set: 174 fall events and 2619 non-fall events (6%)
k-fold testing with varying k to observe lack of data



(a) Decision tree model



(b) Random forest model with 10 decision trees

- ▶ SER: variants give similar or better results
- ▶ STRUT: STRUT_{IG} better, STRUT_{HI} better *only when enough data*

- ▶ Same overall conclusions

Conclusion

Conclusion

- ▶ Good results when comparing with original methods
- ▶ Choosing one method is not easy
- ▶ This suggests a data-dependency of the algorithms

Future works

- ▶ Explore meta models for transfer procedures
- ▶ What if features change ? Consider heterogeneous transfer
- ▶ Methods fitted for decision trees, other model-based transfer might be investigated (i.e. neural networks)
- ▶ If access to source data, then use of domain adaptation techniques

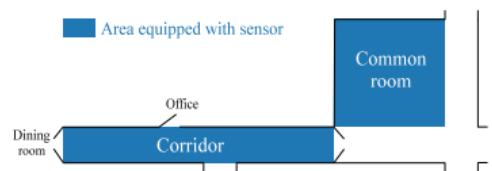
Elderly activity recognition with convolutional neural networks

Introduction

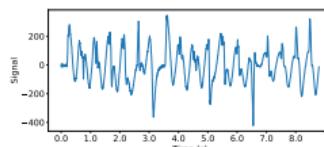
- ▶ **Issue:** one-dimensional signals for large areas
- ▶ **Goal:** Classify elderly from other individuals

Data

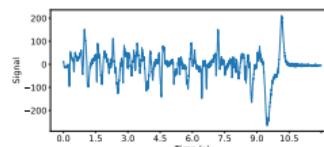
- ▶ 146 signals recorded in a nursing home
- ▶ Walks (one or several persons),
Wheelchair (manual, electrical, pushed
by another), Walk with carts...



Sublabel	Staff	Elderly
Single walk	43	31
Multiple walks	19	11
Other	19	23



Caregiver walk



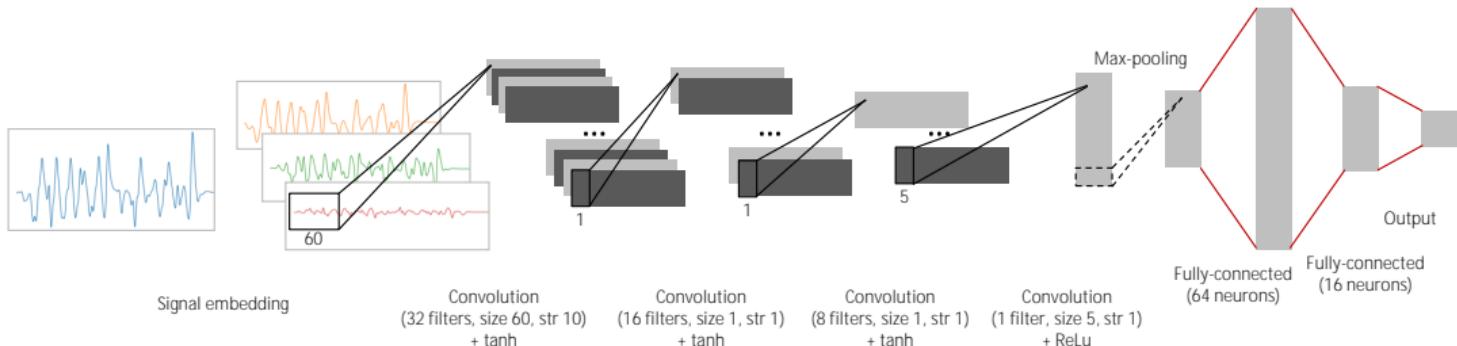
Elderly walk

- ▶ Most signals are made of walks of staff individuals
- ▶ Idea: Bring the model's attention over step-related signals

What we propose: NURSENET , a model to recognize elderly activity using a convolutional neural network and three training steps:

1. Signal embedding using convolutional dictionary learning
2. Step proposal network inspired from Region proposal network
3. The final classification task

NURSENET



- ▶ Use the convolutional representation to “boost” training: First layer (Signal embedding) of NURSENET is trained **separately** using convolutional dictionary learning

Signal embedding

Convolutional dictionary learning

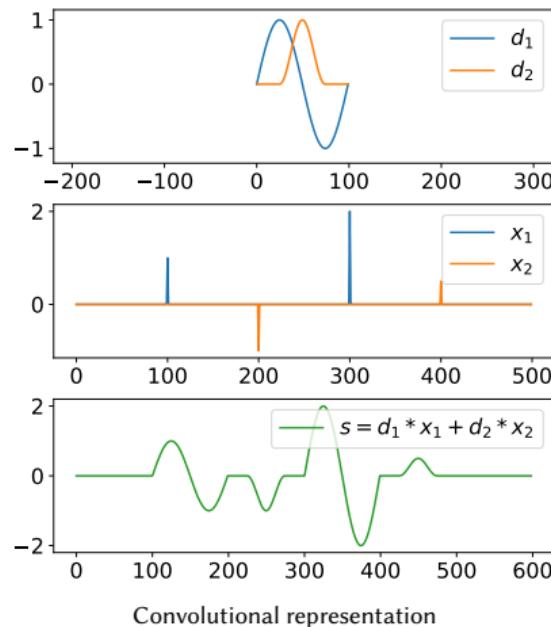
- ▶ \mathbf{s} : data to be represented
- ▶ Objective : find M atoms \mathbf{d}_m and activation signals \mathbf{x}_m such that

$$\mathbf{s} \approx \sum_{m=1}^M \mathbf{x}_m * \mathbf{d}_m$$

- ▶ $*$: convolution

CDL general problem:

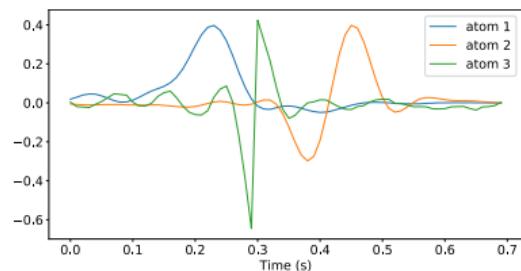
$$\begin{aligned} \arg \min_{\mathbf{x}_m, \mathbf{d}_m} & \frac{1}{2} \left\| \sum_{m=1}^M \mathbf{x}_m * \mathbf{d}_m - \mathbf{s} \right\|_2^2 + \lambda \sum_{m=1}^M \|\mathbf{x}_m\|_1 \\ \text{s.t. } & \|\mathbf{d}_m\|_2 \leq 1 \quad \forall m . \end{aligned}$$



Signal embedding

Learning step atoms

- ▶ Learning with Alternating Direction Method of Multipliers (ADMM) (Bristow et al. [3])
- ▶ 3 atoms of length 0.7 second

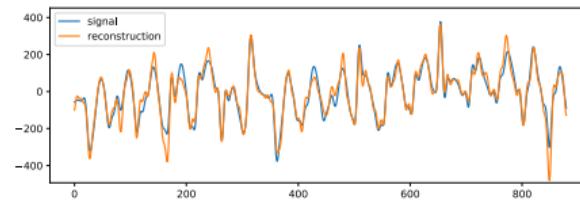
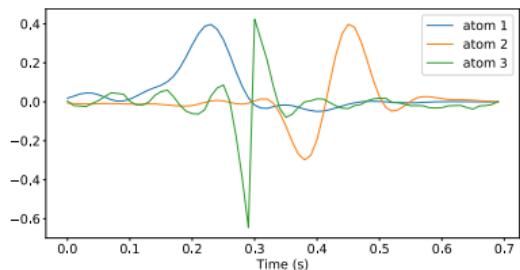


Dictionary

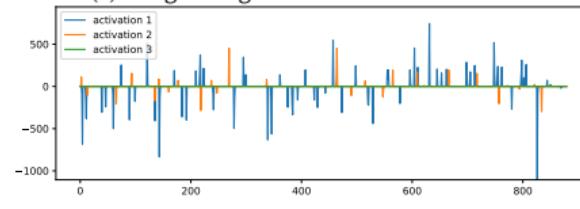
Signal embedding

Learning step atoms

- ▶ Learning with Alternating Direction Method of Multipliers (ADMM) (Bristow et al. [3])
- ▶ 3 atoms of length 0.7 second



(a) Original signal and its reconstruction



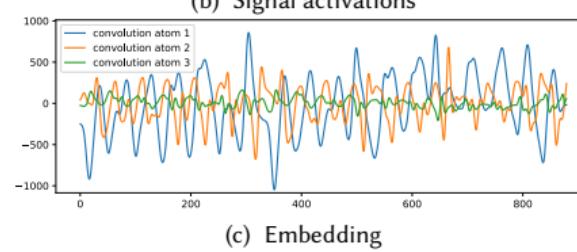
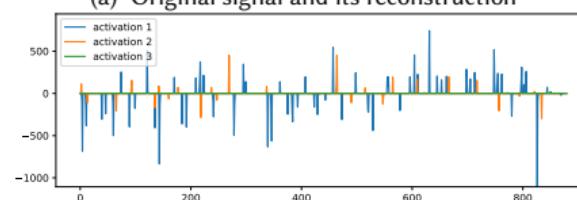
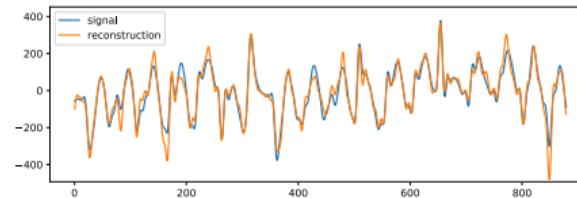
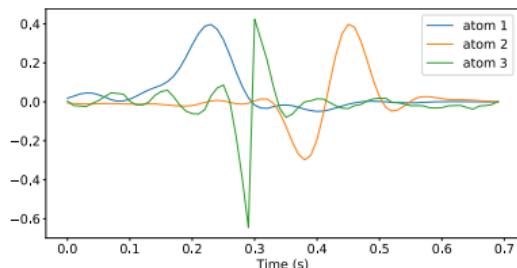
(b) Signal activations

Signal embedding

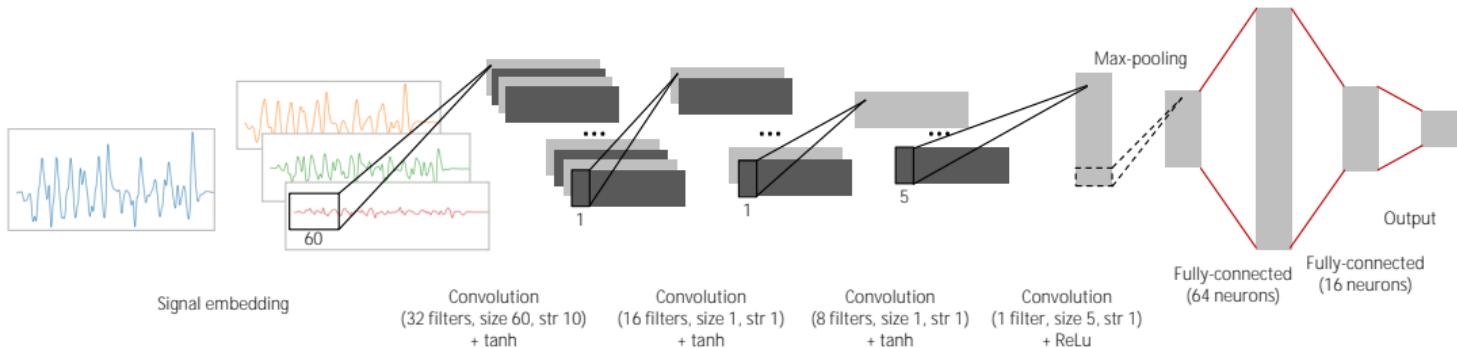
Learning step atoms

- ▶ Learning with Alternating Direction Method of Multipliers (ADMM) (Bristow et al. [3])
- ▶ 3 atoms of length 0.7 second
- ▶ Use the following embedding:

$$\mathbf{S} \doteq (\mathbf{s} * \mathbf{d}_m)_{1 \leq m \leq 3}$$



NURSENET

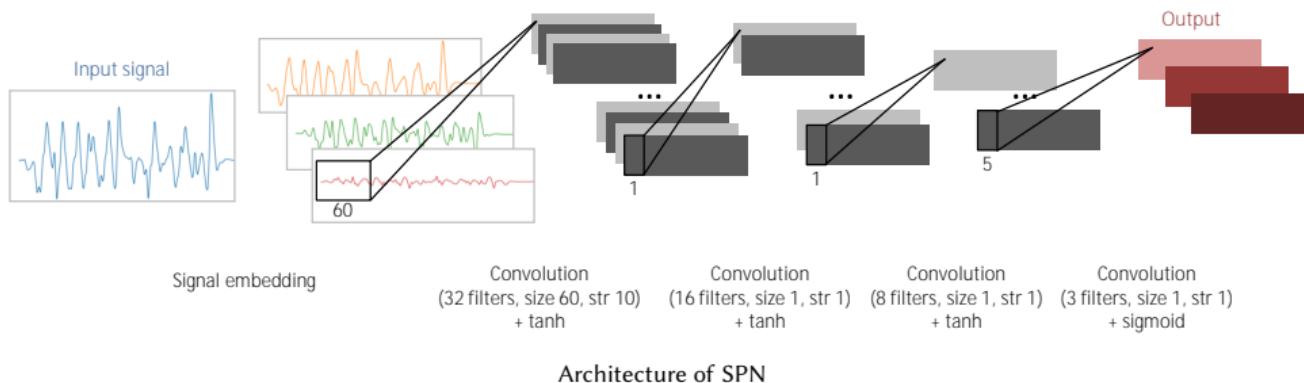


► Second part: Step Proposal Network (SPN)

Step proposal network

Main architecture

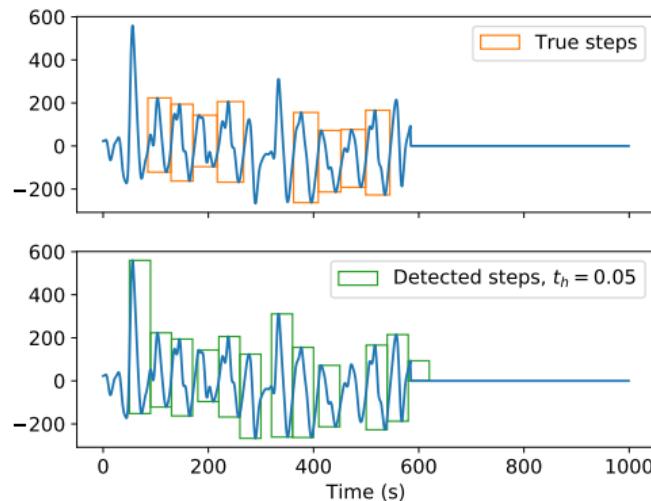
- ▶ Object detection: how to localize objects ?
- ▶ Faster R-CNN (Ren et al. [4]) use the Region proposal network
- ▶ Simple architecture with three hidden layers, all **convolutional**
- ▶ Output: probability of having a step at a specific window location and size
 - ▶ Here 3 sizes and all discrete locations are considered



Step proposal network

Results

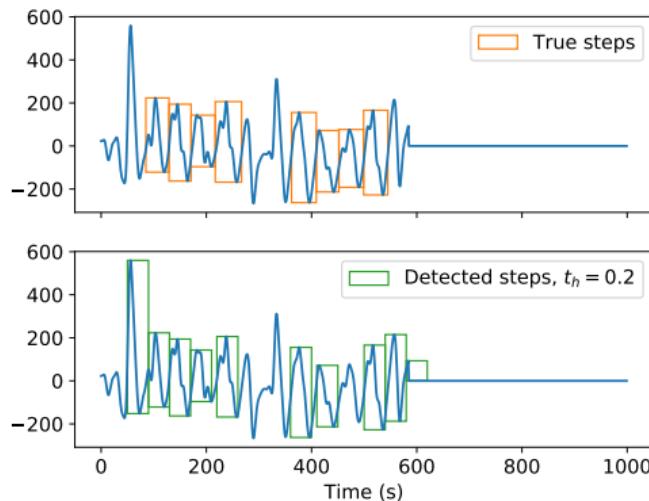
- ▶ Object detection use the mean Average Precision (mAP): area under the Precision-Recall curve
- ▶ mAP = 78,6%



Step proposal network

Results

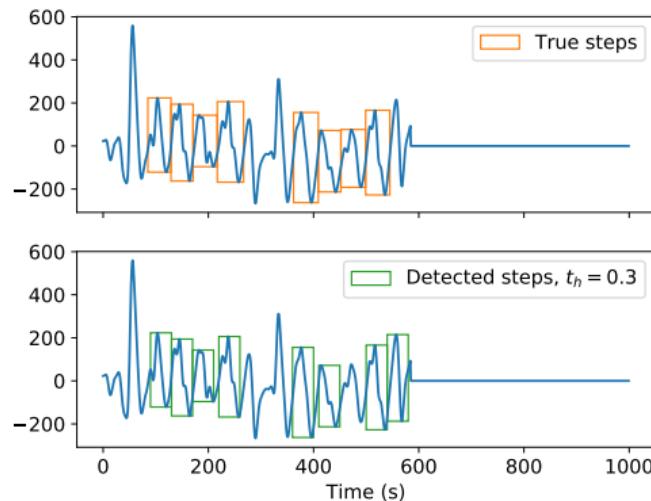
- ▶ Object detection use the mean Average Precision (mAP): area under the Precision-Recall curve
- ▶ mAP = 78,6%



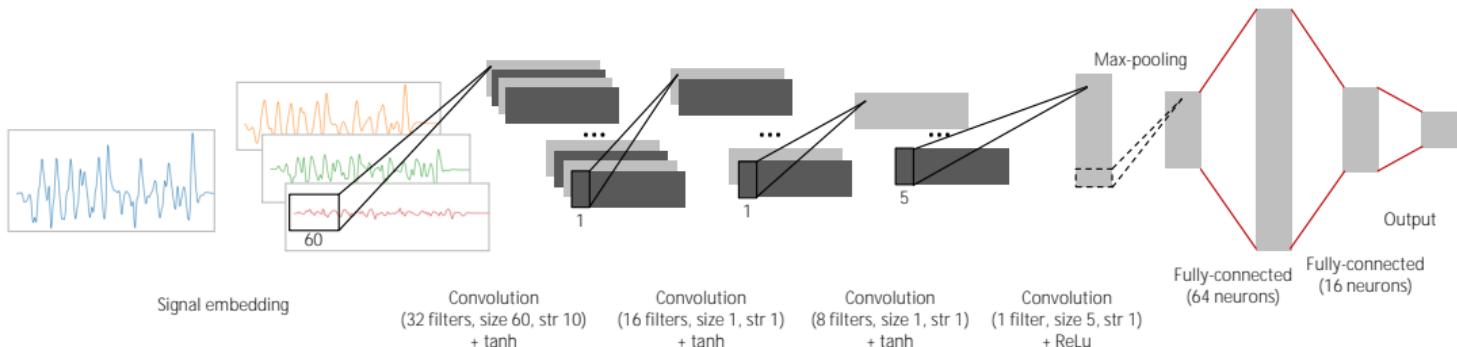
Step proposal network

Results

- ▶ Object detection use the mean Average Precision (mAP): area under the Precision-Recall curve
- ▶ mAP = 78,6%



NURSENET



► Final part: the classification task

Classification

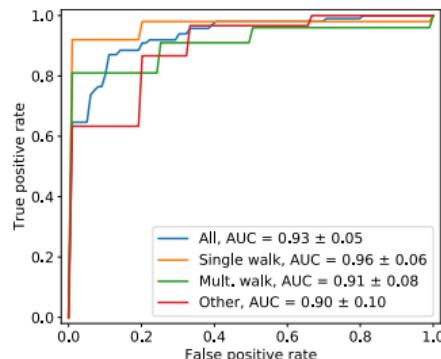
Training

Classification loss

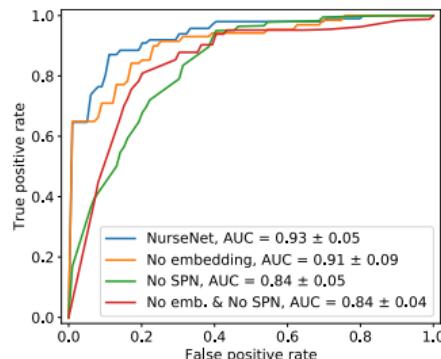
$$\mathcal{L}(\mathbf{s}_i, \hat{y}_i) = -(\mathbb{1}_{y_i=1} \log(\hat{y}_i) + \mathbb{1}_{y_i=0} \log(1 - \hat{y}_i))$$

- ▶ Use the whole set of 146 signals, labeled as *Elderly* and *Staff*
- ▶ NURSENET is trained using classical gradient descent
- ▶ Training time: < 5 minutes
- ▶ Inference (detection over a 10s signal): < 1 second
- ▶ Optimization details
 - ▶ learning rate of 10^{-3}
 - ▶ learning rate decay ($\times 0.9$ every 10 epochs)
 - ▶ Nesterov momentum

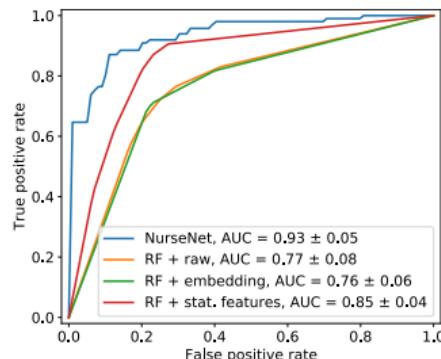
Results



(a) NURSENET results



(b) NURSENET ablation analysis



(c) NURSENET and several RF-based models

Conclusion

- ▶ Good results over all sub-labels
- ▶ Interest of the successive trainings
- ▶ Better results than a RF with statistical features

Future work

- ▶ Improve step detection precision (with regression layer ?)
- ▶ Test on benchmarks step data sets
- ▶ Is it possible to distinguish activities or individuals ?

Conclusion

Conclusion

Contributions



- ▶ A simple and practical model for fall detection
- ▶ Transfer procedure for decision tree adapted to class imbalance
- ▶ A model to distinguish elderly vs. other with high accuracy

Publications

- ▶ L. Minvielle, M. Atiq, R. Serra, M. Mougeot, and N. Vayatis. [Fall detection using smart floor sensor and supervised learning](#). In *2017 39th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, pages 3445–3448, July 2017
- ▶ L. Minvielle, M. Atiq, S. Peignier, and M. Mougeot. [Transfer learning on decision tree with class imbalance](#). In *2019 IEEE 31st International Conference on Tools with Artificial Intelligence (ICTAI)*, pages 1003–1010, Nov 2019
- ▶ L. Minvielle and J. Audiffren. [Nursenet: Monitoring elderly levels of activity with a piezoelectric floor](#). *Sensors*, 19(18), 2019
- ▶ P. Humbert, B. Le Bars, L. Minvielle, and N. Vayatis. [Robust Kernel Density Estimation with Median-of-Means principle](#). *Submitted for publication*, 2020
- ▶ M. Atiq, L. Minvielle, C. Truong. [A data set for fall detection with smart floor sensors](#). *In preparation*, 2020

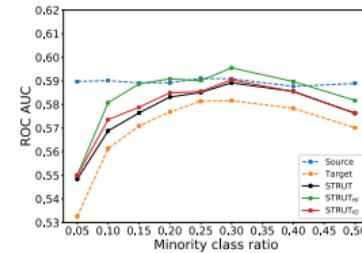
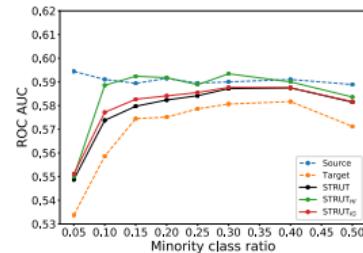
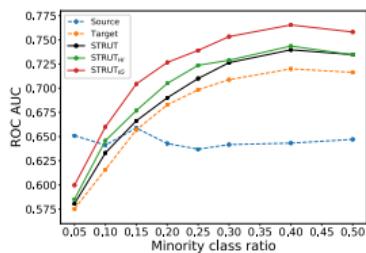
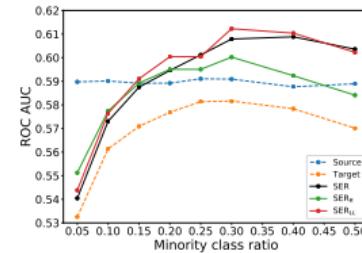
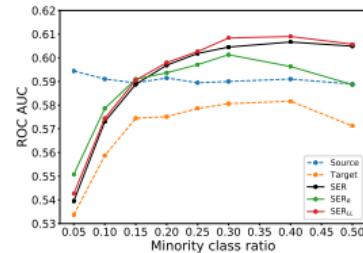
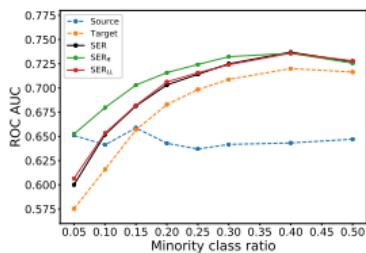
References

- [1] Leo Breiman. Random forests. *Machine Learning*, 45(1):5–32, 2001.
- [2] Leo Breiman, Jerome H. Friedman, Richard A. Olshen, and Charles J. Stone. *Classification and regression trees*. Monterey, CA: Wadsworth & Brooks/Cole Advanced Books & Software, 1984.
- [3] Hilton Bristow, Anders Eriksson, and Simon Lucey. Fast convolutional sparse coding. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2013.
- [4] Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun. Faster r-cnn: Towards real-time object detection with region proposal networks. In *Advances in neural information processing systems*, pages 91–99, 2015.
- [5] N. Segev, M. Harel, S. Mannor, K. Crammer, and R. El-Yaniv. Learn on source, refine on target: A model transfer learning framework with random forests. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 39(9):1811–1824, Sep. 2017.
- [6] World Health Organization. WHO Global Report on Falls Prevention in Older Age. *Community Health*, page 53, 2007. URL https://www.who.int/ageing/publications/Falls_prevention7March.pdf.

Results

Public data

Two real data sets where the imbalance ratio is controlled with downsampling.



(a) Magic Gamma Telescope

(b) Office-Caltech: *amazon* → *webcam*

(c) Office-Caltech: *caltech* → *webcam*

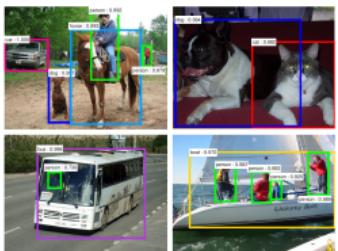
- ▶ SER: SER_{LL} close to SER, SER_R either better... or worse
- ▶ STRUT: both variants perform as well or better

Step proposal network

Region proposal network

Object detection

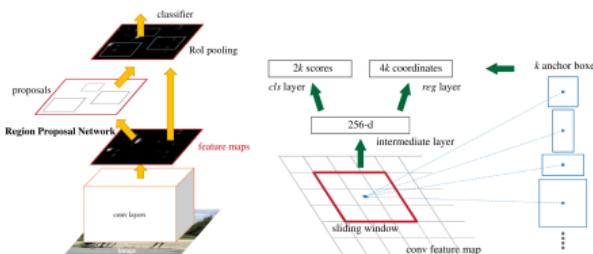
- ▶ Object detection: Where are the objects and what are they classes ?
- ▶ How to efficiently localize objects ?
- ▶ Proposal models
- ▶ Faster R-CNN (Ren et al. [4]) use the Region proposal network



Object detection. Source:Ren et al. [4]

Region proposal network

- ▶ Main idea: proposals are generated by a CNN
- ▶ A sliding window is passed: multiple *anchors* over each location (various sizes and scales)
- ▶ Two layers: Classification (Object / Not Object) and Regression (anchor coordinates)



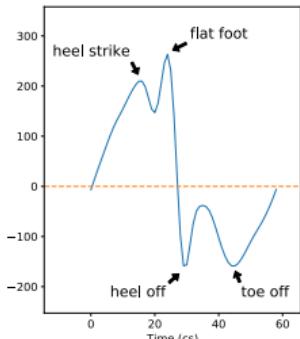
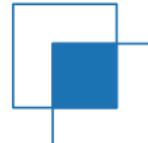
Region proposal network. Source: Ren et al. [4]

Step proposal network

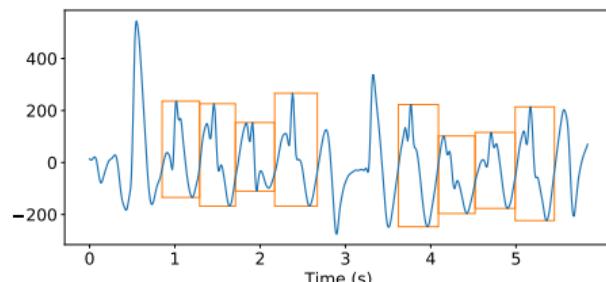
Principle

- ▶ Objective of SPN : output boxes with largest Intersection over Union (IoU)
- ▶ IoU: \mathbf{b}_j are labelled boxes, \hat{b} is an estimated box:

$$\text{IoU}(\hat{b}) \doteq \max_j \frac{|\mathbf{b}_j \cap \hat{b}|}{|\mathbf{b}_j \cup \hat{b}|}$$



(a) Step signal



(b) Step labels over a walk signal

Step proposal network

Training

Boxes and loss

- ▶ Output: a matrix $\mathbf{W} \in \mathbb{R}^{T \times K}$
 - ▶ T : signal length
 - ▶ K : number of different box sizes
- ▶ $\mathbf{W}_{t,k}$: probability that the box b_t^k starting at time t and of size 0.4s, 0.5s, or 0.6s (for respectively $k = 1, 2$, or 3) has a large IoU score
- ▶ Positive boxes: $\text{IoU}(b_t^k) > \sqrt{0.7}$
- ▶ Negative boxes: $\text{IoU}(b_t^k) < \sqrt{0.3}$
- ▶ Other are not used for training

The loss function \mathcal{L} over a signal \mathbf{s} is defined as:

$$\mathcal{L}(\mathbf{s}, \mathbf{W}) = \sum_t \sum_{k \in [1, 2, 3]} \mathbb{1}_{\text{IoU}(b_t^k) > \sqrt{0.7}} \log(\mathbf{W}_{t,k}) + \mathbb{1}_{\text{IoU}(b_t^k) < \sqrt{0.3}} \log(1 - \mathbf{W}_{t,k}).$$

Procedure

- ▶ Use the 43 walk signals from caregivers
- ▶ SPN is trained using classical gradient descent
- ▶ Training time: < 5 minutes
- ▶ Inference (detection over a 10s signal): < 1 second
- ▶ Optimization details
 - ▶ learning rate of 10^{-3}
 - ▶ learning rate decay ($\times 0.9$ every 10 epochs)
 - ▶ Nesterov momentum