

Détection d'événements à partir de capteurs sols – application au suivi de personnes fragiles

Soutenance de thèse

Ludovic Minvielle

Thèse industrielle entre le Centre Borelli (ENS Paris-Saclay) et Tarkett

Mercredi 15 Juillet 2020



Introduction

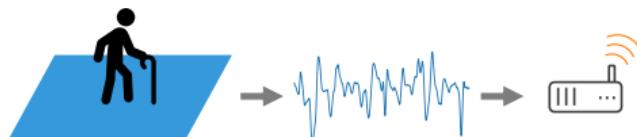
General context

- ▶ Elderly population is growing
- ▶ Higher levels of frailty globally
- ▶ Increasing demand for reliable monitoring devices
- ▶ Tarkett, French company: 12,500 employees, 13 industrial sites, 1.3 millions m² of flooring per day
- ▶ *Floor in Motion*: a floor-based sensor for elderly care
- ▶ **Objective**: providing tools for elderly monitoring in nursing homes

Introduction

General context

- ▶ Elderly population is growing
- ▶ Higher levels of frailty globally
- ▶ Increasing demand for reliable monitoring devices
- ▶ Tarkett, French company: 12,500 employees, 13 industrial sites, 1.3 millions m² of flooring per day
- ▶ *Floor in Motion*: a floor-based sensor for elderly care
- ▶ **Objective:** providing tools for elderly monitoring in nursing homes



Introduction

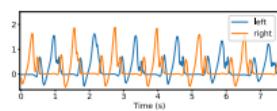
General context

- ▶ Elderly population is growing
- ▶ Higher levels of frailty globally
- ▶ Increasing demand for reliable monitoring devices
- ▶ Tarkett, French company: 12,500 employees, 13 industrial sites, 1.3 millions m² of flooring per day
- ▶ *Floor in Motion*: a floor-based sensor for elderly care
- ▶ **Objective:** providing tools for elderly monitoring in nursing homes

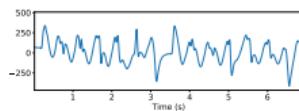


Scientific context

- ▶ Processing and understanding time series
- ▶ Real world application



Foot-attached
accelerometer



Tarkett's floor sensor

Introduction

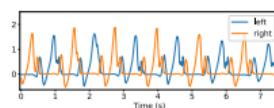
General context

- ▶ Elderly population is growing
- ▶ Higher levels of frailty globally
- ▶ Increasing demand for reliable monitoring devices
- ▶ Tarkett, French company: 12,500 employees, 13 industrial sites, 1.3 millions m² of flooring per day
- ▶ *Floor in Motion*: a floor-based sensor for elderly care
- ▶ **Objective:** providing tools for elderly monitoring in nursing homes

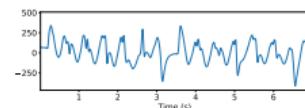


Scientific context

- ▶ Processing and understanding time series
- ▶ Real world application



Foot-attached
accelerometer



Tarkett's floor sensor

Questions

- ▶ Is a floor-based system a good idea ?
- ▶ How to build a fall detector out of it ?
- ▶ How to use real data to improve our model ?
- ▶ The signal is unidimensional: can we still distinguish people ?

Table of contents

1. Monitoring systems
2. Fall detection
3. Transfer learning
4. Elderly activity recognition
5. Conclusion

Monitoring systems for fall detection

Sensors

What makes a good monitoring system ?

- ▶ coverage and occlusion
- ▶ intrusiveness
- ▶ signal quality / information
- ▶ robustness
- ▶ ease of installation / use
- ▶ scalability

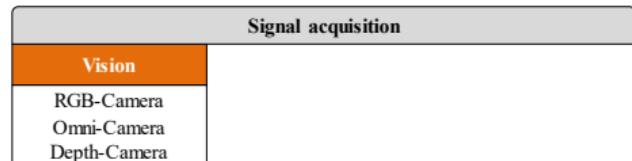
Criteria

Coverage/Occlusion
Intrusiveness
Signal quality / info
Robustness
Ease of instal. / use
Scalability

Sensors

What makes a good monitoring system ?

- ▶ coverage and occlusion
- ▶ intrusiveness
- ▶ signal quality / information
- ▶ robustness
- ▶ ease of installation / use
- ▶ scalability



Criteria	RGB cam	Depth cam
Coverage/Occlusion	★☆☆	★☆☆
Intrusiveness	★☆☆	★☆☆
Signal quality / info	★★★	★★★
Robustness	★☆☆	★★★
Ease of instal. / use	★☆☆	★☆☆
Scalability	★☆☆	★☆☆

Sensors

What makes a good monitoring system ?

- ▶ coverage and occlusion
- ▶ intrusiveness
- ▶ signal quality / information
- ▶ robustness
- ▶ ease of installation / use
- ▶ scalability

Signal acquisition	
Vision	Wearable
RGB-Camera	
Omni-Camera	Accelerometer
Depth-Camera	Gyroscope
	Barometric pressure

Criteria	RGB cam	Depth cam	Wearable
Coverage/Occlusion	★☆☆	★☆☆	★★★
Intrusiveness	★☆☆	★☆☆	★★☆
Signal quality / info	★★★	★★★	★★☆
Robustness	★☆☆	★★★	★★★
Ease of instal. / use	★☆☆	★☆☆	★★☆
Scalability	★☆☆	★☆☆	★★★

Sensors

What makes a good monitoring system ?

- ▶ coverage and occlusion
- ▶ intrusiveness
- ▶ signal quality / information
- ▶ robustness
- ▶ ease of installation / use
- ▶ scalability

Signal acquisition		
Vision	Wearable	Ambient
RGB-Camera	Accelerometer	Microphone
Omni-Camera	Gyroscope	Radar
Depth-Camera	Barometric pressure	Wi-Fi
		Vibration
		Pressure

Criteria	RGB cam	Depth cam	Wearable	Acoustic	Radar / Wi-Fi	Vibration	Floor
Coverage/Occlusion	★☆☆	★☆☆	★★★	★☆☆	★☆☆	★★★	★★★
Intrusiveness	★☆☆	★☆☆	★☆☆	★☆☆	★☆☆	★★★	★★★
Signal quality / info	★★★	★★★	★☆☆	★☆☆	★☆☆	★☆☆	★☆☆
Robustness	★☆☆	★★★	★★★	★☆☆	★☆☆	★☆☆	★☆☆
Ease of instal. / use	★☆☆	★☆☆	★☆☆	★☆☆	★☆☆	★★★	★☆☆
Scalability	★☆☆	★☆☆	★★★	★☆☆	★☆☆	★☆☆	★★★

Information extraction

How to process the inputs ?

- ▶ All systems use feature extraction
- ▶ The “level” of feature engineering depends on the complexity / dimensionality of the input signal

How to deal with processed signals ?

Time series classification

1. Series as *sequences*
 - ▶ Distance-based methods
2. Series as *feature vectors*
 - ▶ Computing several measures over a fixed size
 - ▶ Classification models (Anomaly detection, classical supervised models...)

Signal acquisition		
Vision	Wearable	Ambient
RGB-Camera	Accelerometer	Microphone
Omni-Camera	Gyroscope	Radar
Depth-Camera	Barometric pressure	Wi-Fi
		Vibrational Pressure



Feature extraction		
Vision	Wearable	Ambient
Position	Position	Statistical measures
Motion	Velocity	Fourier transform
Shape	Angle	Wavelet transform
		Cepstrum features



Decision rule	
Threshold	Machine learning
	kNN SVM HMM Decision Tree

A fall detection system

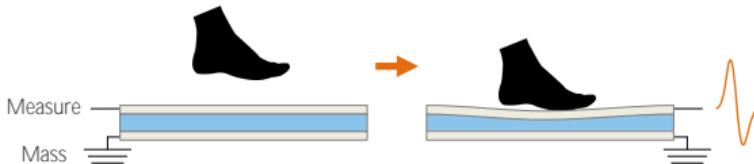
Tarkett sensor

- ▶ Piezoelectric principle:

$$d = \frac{Q}{F},$$

(simple version) with d the *piezoelectric constant*.

When stressed or squeezed, the material emits charges.



Tarkett sensor

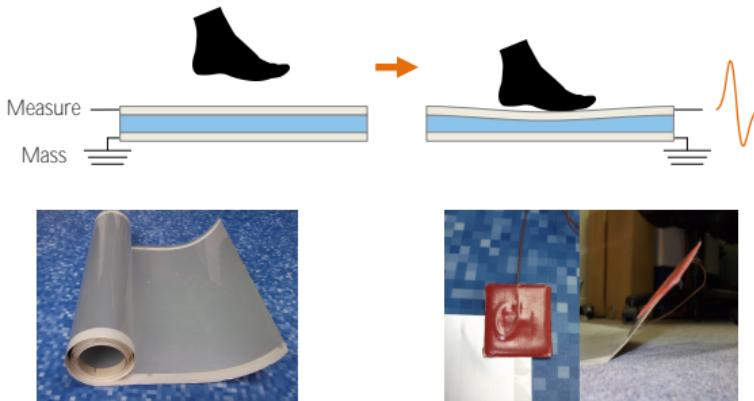
- ▶ Piezoelectric principle:

$$d = \frac{Q}{F},$$

(simple version) with d the *piezoelectric constant*.

When stressed or squeezed, the material emits charges.

- ▶ How does this look like ?
0.3 mm thick and 60 cm wide roll
with customizable length



Tarkett sensor

- ▶ Piezoelectric principle:

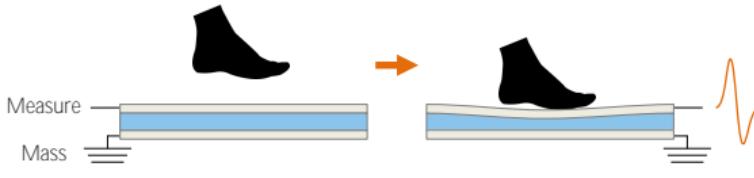
$$d = \frac{Q}{F},$$

(simple version) with d the *piezoelectric constant*.

When stressed or squeezed, the material emits charges.

- ▶ How does this look like ?
0.3 mm thick and 60 cm wide roll
with customizable length
- ▶ How is it installed ?

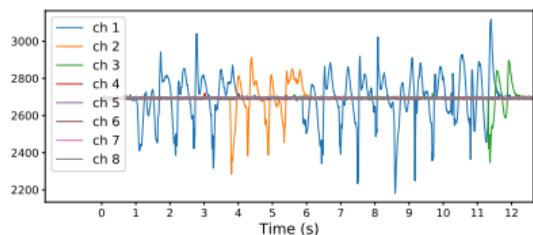
- ▶ Under the flooring
- ▶ Several connected bands for each area, hence one area corresponds to one input



Data

Preprocessing

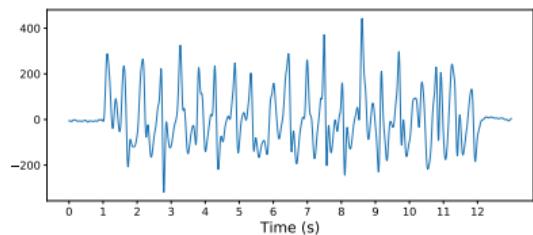
- ▶ linear detrending
- ▶ low-pass filtering
- ▶ zeroing low energy channels
- ▶ sum over all channels



Data

Preprocessing

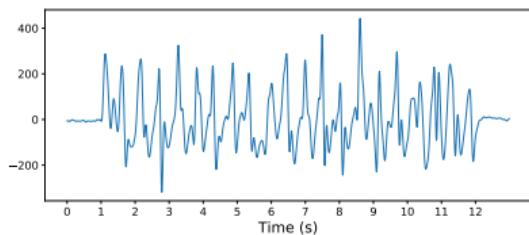
- ▶ linear detrending
- ▶ low-pass filtering
- ▶ zeroing low energy channels
- ▶ sum over all channels



Data

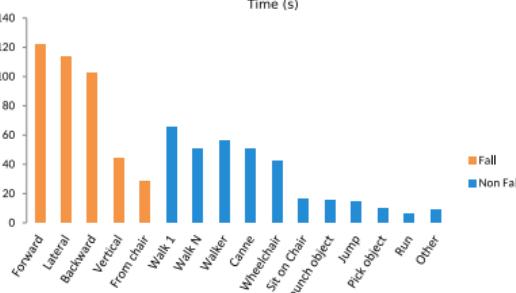
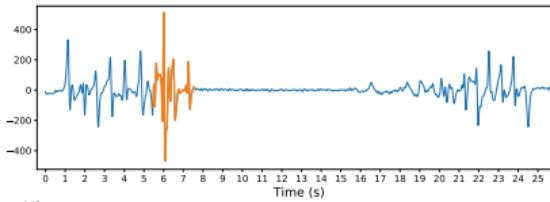
Preprocessing

- ▶ linear detrending
- ▶ low-pass filtering
- ▶ zeroing low energy channels
- ▶ sum over all channels



Experimental dataset

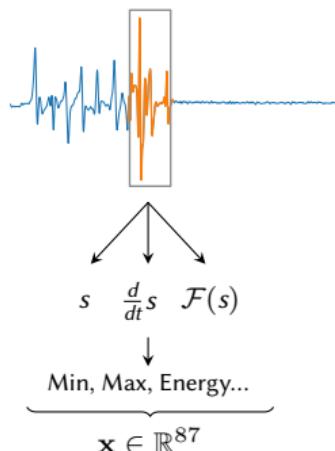
- ▶ 742 signals
- ▶ 55% fall, 45% non-fall
- ▶ varied fall events (forward, backward...) and activities of daily living (walking, sitting...)



Method

Time series as *feature vector*. At every timestamp:

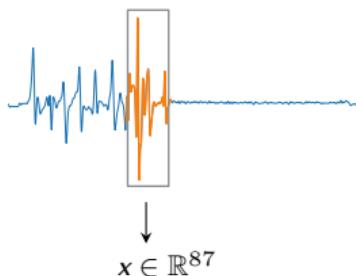
1. Window over the signal: 2.5 s
2. Compute feature vector: 29 statistical measures
(Min, Max, Shannon energy, Percentile,...) over
three representations of the signal



Method

Time series as *feature vector*. At every timestamp:

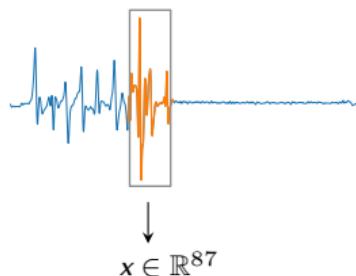
1. Window over the signal: 2.5 s
2. Compute feature vector: 29 statistical measures
(Min, Max, Shannon energy, Percentile,...) over
three representations of the signal



Method

Time series as *feature vector*. At every timestamp:

1. Window over the signal: 2.5 s
2. Compute feature vector: 29 statistical measures (Min, Max, Shannon energy, Percentile,...) over three representations of the signal



3. Classification model: Random Forest [1], based on **decision trees**

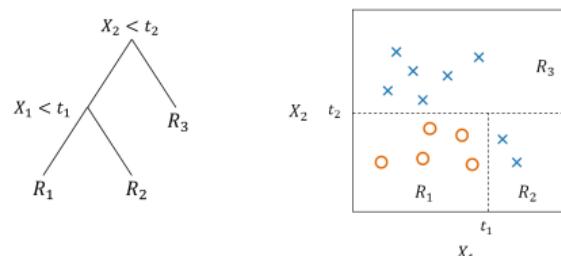
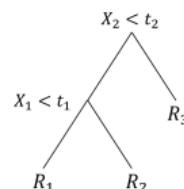
Decision tree

Feature space $\mathcal{X} = \mathbb{R}^Q$. Division of \mathcal{X} into non-overlapping regions R_1, \dots, R_j . Algorithm CART: recursive binary splits [2] that solve:

$$\arg \min_{X_q, \tau} \text{IG} ,$$

$$\text{with } \text{IG}(X_q, \tau) = I(n) - \frac{N_l}{N_n} I(l) - \frac{N_r}{N_n} I(r) ,$$

$$\text{and } I(n) = \text{Gini}(n) = \sum_k p_{nk}(1 - p_{nk}) .$$



$$\text{Prediction function: } f(x) = \sum_{j=1}^J c_j \mathbb{1}(x \in R_j)$$

Method

Random forest

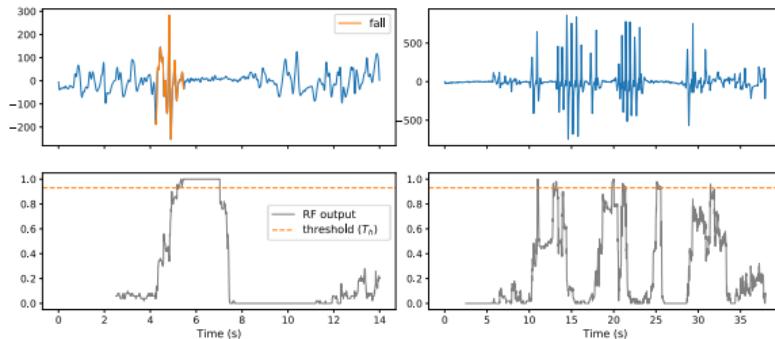
Decision trees d_1, \dots, d_{N_T} grown with two rules:

- ▶ Each tree is trained with a *bootstrap* of the training set
- ▶ At each split, access to a random subset of pool of features

Each tree is a “vote” for a class. The prediction function is then

$$f(x) = \arg \max_k f_k(x),$$

$$\text{with } f_k(x) = \frac{1}{N_T} \sum_{i=1}^{N_T} \mathbb{1}(d_i(x) = k)$$



Time aggregation

$N_f(t)$: number of trees voting for *fall*

Use a buffer $B_s \in \mathbb{N}$ and a threshold $T_h \in [0, 1]$

$$g(t) = \frac{\sum_{u=t-B_s+1}^t N_f(u)}{B_s \times N_T}$$

New binary classification function: $d(t) = \begin{cases} 1, & \text{if } g(t) > T_h \\ 0, & \text{otherwise} \end{cases}$

Method

Random forest

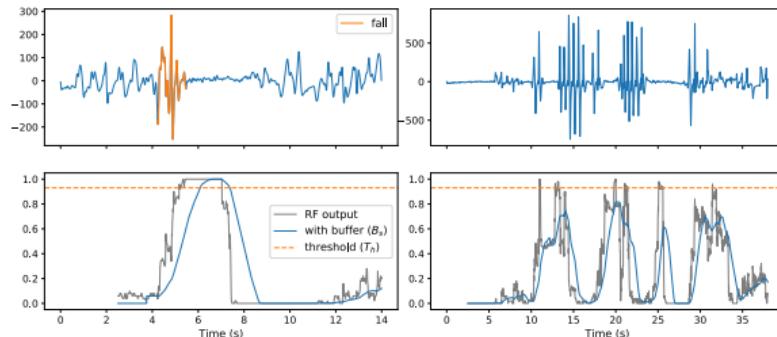
Decision trees d_1, \dots, d_{N_T} grown with two rules:

- ▶ Each tree is trained with a *bootstrap* of the training set
- ▶ At each split, access to a random subset of pool of features

Each tree is a “vote” for a class. The prediction function is then

$$f(x) = \arg \max_k f_k(x),$$

$$\text{with } f_k(x) = \frac{1}{N_T} \sum_{i=1}^{N_T} \mathbb{1}(d_i(x) = k)$$



Time aggregation

$N_f(t)$: number of trees voting for *fall*

Use a buffer $B_s \in \mathbb{N}$ and a threshold $T_h \in [0, 1]$

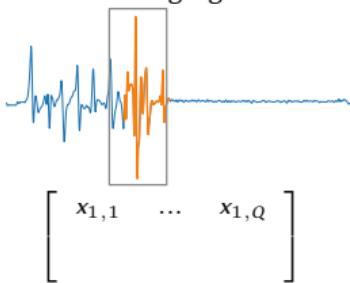
$$g(t) = \frac{\sum_{u=t-B_s+1}^t N_f(u)}{B_s \times N_T}$$

New binary classification function: $d(t) = \begin{cases} 1, & \text{if } g(t) > T_h \\ 0, & \text{otherwise} \end{cases}$

Method

Data augmentation

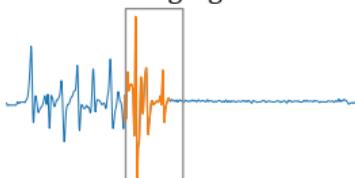
Select r windows in training signals



Method

Data augmentation

Select r windows in training signals

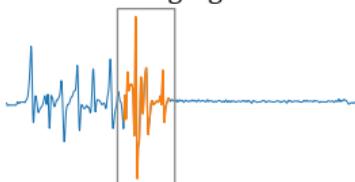


$$\begin{bmatrix} x_{1,1} & \dots & x_{1,Q} \\ x_{2,1} & \dots & x_{2,Q} \end{bmatrix}$$

Method

Data augmentation

Select r windows in training signals

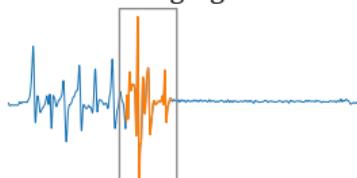


$$\begin{bmatrix} x_{1,1} & \dots & x_{1,Q} \\ x_{2,1} & \dots & x_{2,Q} \\ x_{3,1} & \dots & x_{3,Q} \end{bmatrix}$$

Method

Data augmentation

Select r windows in training signals

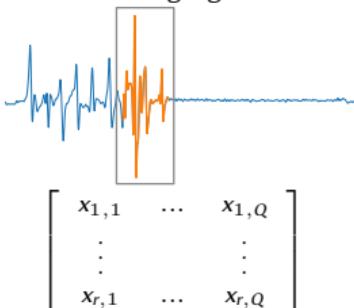


$$\begin{bmatrix} x_{1,1} & \dots & x_{1,Q} \\ \vdots & & \vdots \\ x_{r,1} & \dots & x_{r,Q} \end{bmatrix}$$

Method

Data augmentation

Select r windows in training signals

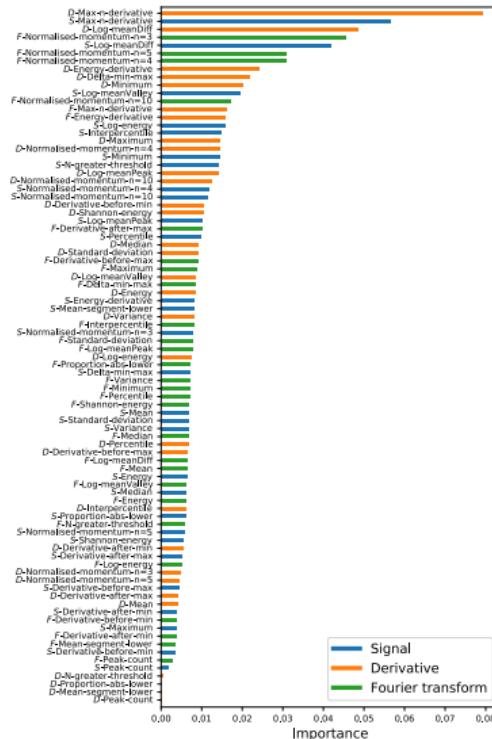


Feature reduction

Feature importance

$$\text{Tree: } I(X_q) = \sum_{\text{nodes } t} p(t) \Delta i(t) \mathbb{1}(v(t) = X_q)$$

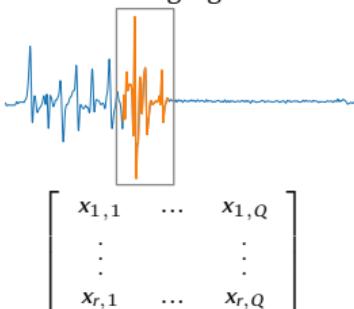
$$\text{Random forest: } I(X_q) = \frac{1}{N_T} \sum_{n=1}^{N_T} I(T_n, X_q)$$



Method

Data augmentation

Select r windows in training signals



Feature reduction

Feature importance

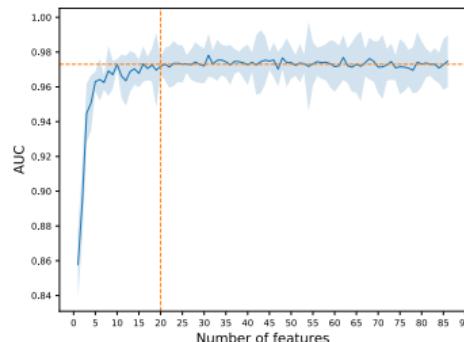
$$\text{Tree: } I(X_q) = \sum_{\text{nodes } t} p(t) \Delta i(t) \mathbb{1}(v(t) = X_q)$$

$$\text{Random forest: } I(X_q) = \frac{1}{N_T} \sum_{n=1}^{N_T} I(T_n, X_q)$$

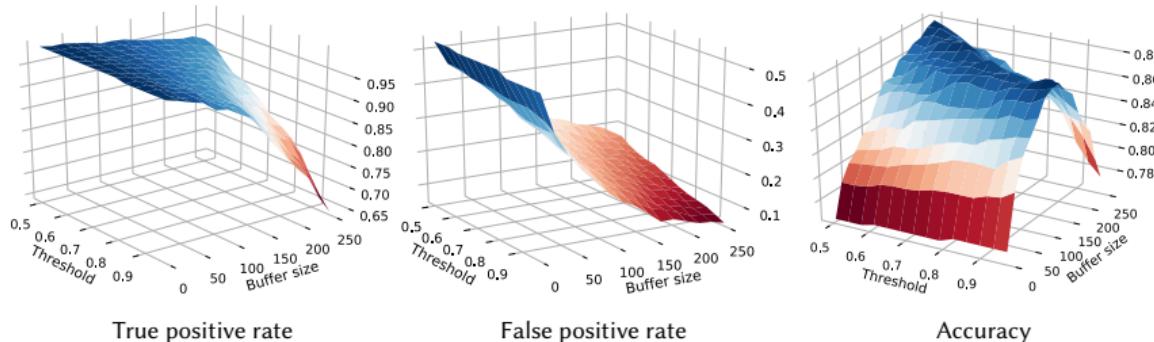
Recursive feature elimination

Initial pool of Q features X_1, \dots, X_Q .

1. Train several times and record variable importances
2. Average of importances over trainings.
 $X_{q*} = \arg \min_{X_i} I(X_i)$
3. Remove X_{q*} from the pool of features and back to step 1



Results



True positive rate

False positive rate

Accuracy

Model	Accuracy	TPR	FPR	$\text{TPR}_{\text{FPR} < 10}^{\text{min}}$	$\text{TPR}_{\text{FPR} < 10}^{\text{max}}$
$r = 5, Q = 20$					
LR	86.8 ± 1.5	90.5 ± 2.4	17.7 ± 4.9	67.0 ± 10.8	80.4 ± 6.4
LDA	85.5 ± 1.2	91.0 ± 2.1	21.7 ± 3.7	56.9 ± 7.0	78.7 ± 3.8
k-NN	87.0 ± 1.9	89.2 ± 1.4	16.0 ± 4.7	63.1 ± 4.2	83.1 ± 2.5
SVM	87.6 ± 3.2	90.0 ± 4.5	15.5 ± 6.8	69.2 ± 2.1	82.9 ± 3.2
MLP	88.2 ± 1.5	92.4 ± 1.2	17.3 ± 4.1	71.4 ± 4.5	85.1 ± 2.1
RF	88.2 ± 1.5	91.7 ± 3.5	16.2 ± 6.2	63.8 ± 6.8	84.3 ± 7.9

Transfer learning on decision tree

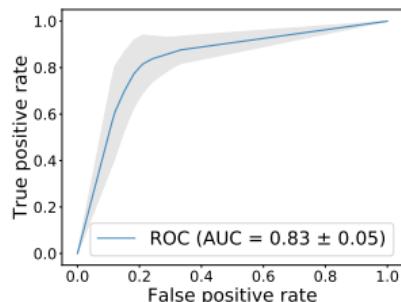
Experimental vs. operational

Transfer learning

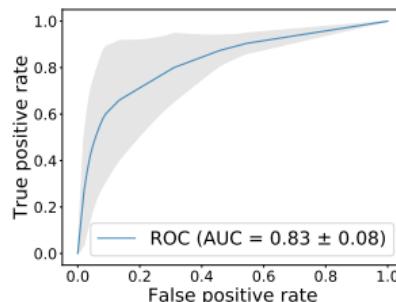
- ▶ Source domain: $\mathcal{D}_S = \{\mathcal{X}_S, P(X_S)\}$
- ▶ Target domain: $\mathcal{D}_T = \{\mathcal{X}_T, P(X_T)\}$
- ▶ Source task: $\mathcal{T}_S = \{\mathcal{Y}_S, f^S\}$
- ▶ Target task: $\mathcal{T}_T = \{\mathcal{Y}_T, f^T\}$

Our case

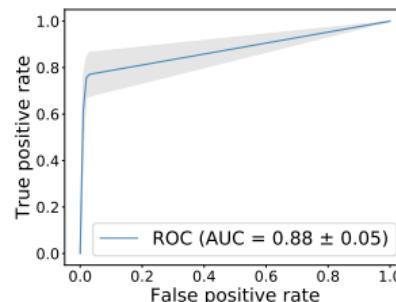
- ▶ $\mathcal{X}_S = \mathcal{X}_T$
- ▶ $P(X_S) \neq P(X_T)$
- ▶ $\mathcal{Y}_S = \mathcal{Y}_T$
- ▶ $f^S \neq f^T$



Source tested on source



Source tested on target

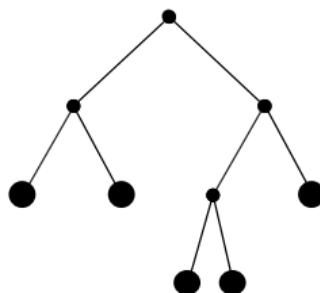


Target tested on target

Model-based transfer

Segev et al. [7]

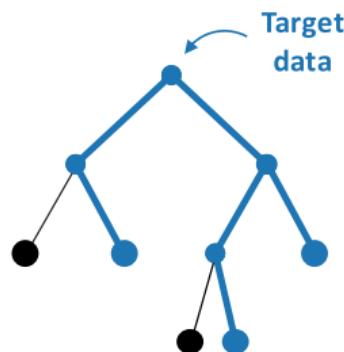
Structure Expansion / Reduction (SER)



Model-based transfer

Segev et al. [7]

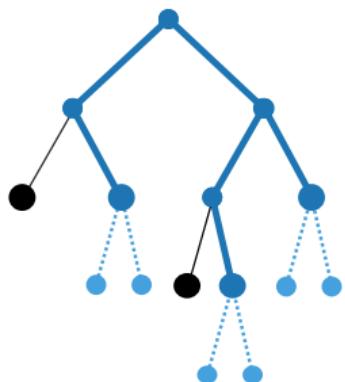
Structure Expansion / Reduction (SER)



Model-based transfer

Segev et al. [7]

Structure Expansion / Reduction (SER)

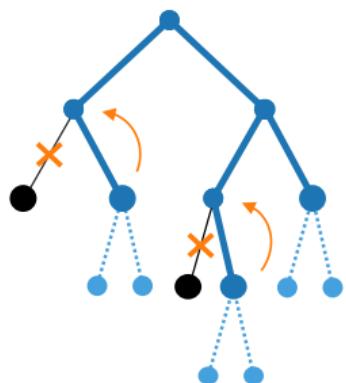


1. Expansion

Model-based transfer

Segev et al. [7]

Structure Expansion / Reduction (SER)

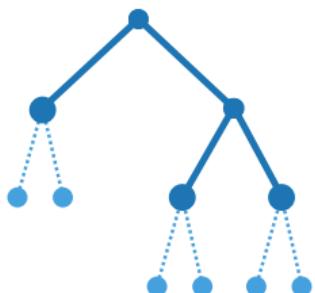


1. Expansion
2. Reduction

Model-based transfer

Segev et al. [7]

Structure Expansion / Reduction (SER)

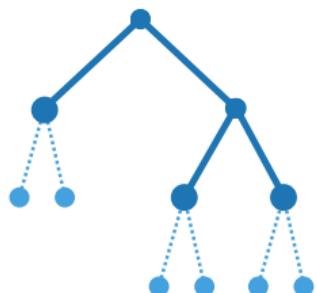


1. Expansion
2. Reduction

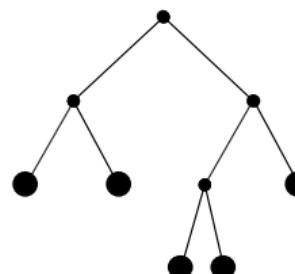
Model-based transfer

Segev et al. [7]

Structure Expansion / Reduction (SER)



Structure Transfer (STRUT)

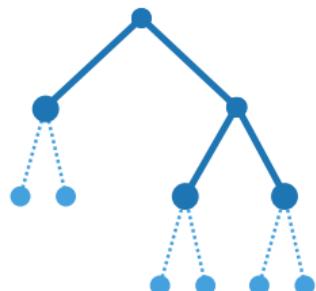


1. Expansion
2. Reduction

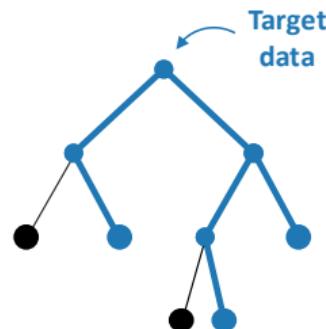
Model-based transfer

Segev et al. [7]

Structure Expansion / Reduction (SER)



Structure Transfer (STRUT)

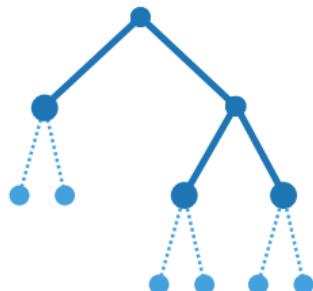


1. Expansion
2. Reduction

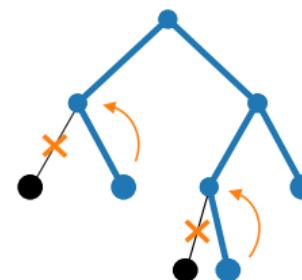
Model-based transfer

Segev et al. [7]

Structure Expansion / Reduction (SER)



Structure Transfer (STRUT)



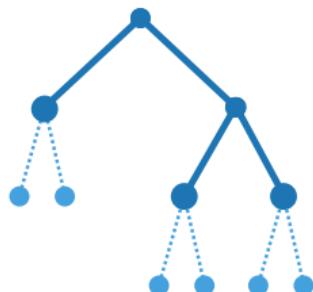
1. Pruning

1. Expansion
2. Reduction

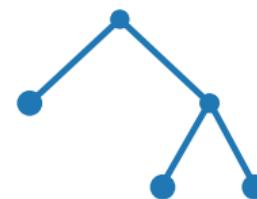
Model-based transfer

Segev et al. [7]

Structure Expansion / Reduction (SER)



Structure Transfer (STRUT)



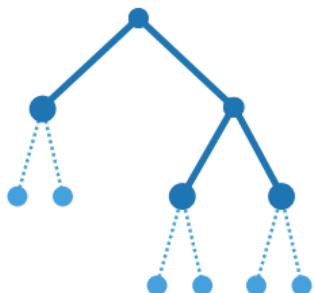
1. Pruning

1. Expansion
2. Reduction

Model-based transfer

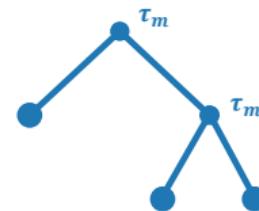
Segev et al. [7]

Structure Expansion / Reduction (SER)



1. Expansion
2. Reduction

Structure Transfer (STRUT)

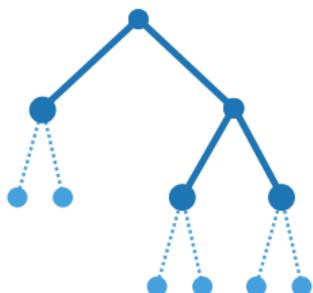


1. Pruning
2. Threshold update

Model-based transfer

Segev et al. [7]

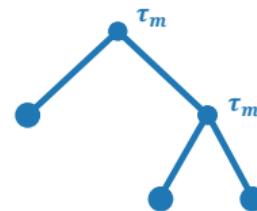
Structure Expansion / Reduction (SER)



1. Expansion
2. Reduction

Partition refinement or simplification

Structure Transfer (STRUT)



1. Pruning
2. Threshold update

Drifts

Leaf loss risk

Homogeneous class imbalance

$$P^T(x|y) = P^S(x|y)$$

$$P^T(y|x) = \lambda_y \frac{P^S(y|x)}{\int \lambda_y P^S(y|x) dy}$$

$$\text{with } \lambda_y = \frac{P^T(y)}{P^S(y)}$$

Leaf loss risk

Homogeneous class imbalance

$$P^T(x|y) = P^S(x|y)$$

$$P^T(y|x) = \lambda_y \frac{P^S(y|x)}{\int \lambda_y P^S(y|x) dy}$$

with $\lambda_y = \frac{P^T(y)}{P^S(y)}$

Leaf loss risk

Significant leaf: Leaf l that conserves the minority class k_{min} after Target update:

$$\forall k \neq k_{min}, \quad P^T(y = k_{min} | x \in l) > P^T(y = k | x \in l)$$

Leaf loss risk:

$$R_L(l) = P^T(x \notin l | y = k_{min})^{n_{k_{min}}}$$

Leaf loss risk

Homogeneous class imbalance

$$P^T(x|y) = P^S(x|y)$$

$$P^T(y|x) = \lambda_y \frac{P^S(y|x)}{\int \lambda_y P^S(y|x) dy}$$

with $\lambda_y = \frac{P^T(y)}{P^S(y)}$

Leaf loss risk

Significant leaf: Leaf l that conserves the minority class k_{min} after Target update:

$$\forall k \neq k_{min}, \quad P^T(y = k_{min} | x \in l) > P^T(y = k | x \in l)$$

Leaf loss risk:

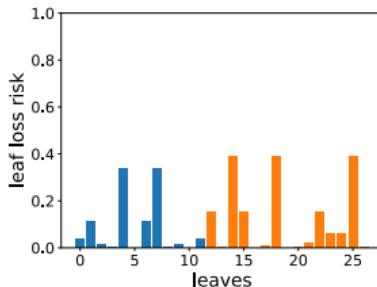
$$R_L(l) = P^T(x \notin l | y = k_{min})^{n_{k_{min}}}$$

Leaf loss risk under homogeneous class imbalance

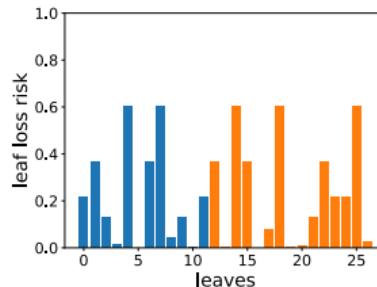
$$\forall k \neq k_{min}, \quad \lambda_{k_{min}} P^S(y = k_{min} | x \in l) > \lambda_k P^S(y = k | x \in l)$$

$$R_L(l) = P^S(x \notin l | y = k_{min})^{n_{k_{min}}}$$

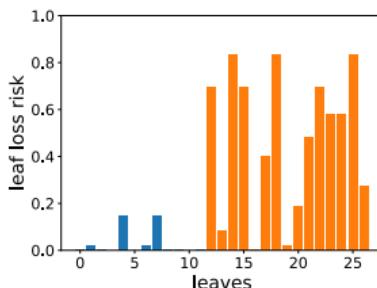
Leaf loss risk



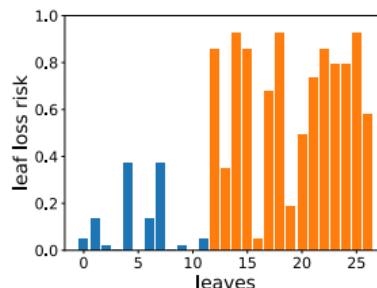
(a) Balanced data with 200 instances



(b) Balanced data with 100 instances



(c) Imbalanced data (10% ratio) with 200 instances

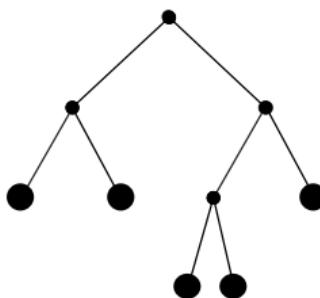


(d) Imbalanced data (10% ratio) with 100 instances

SER for class imbalance

SER_R, SER_{LL}

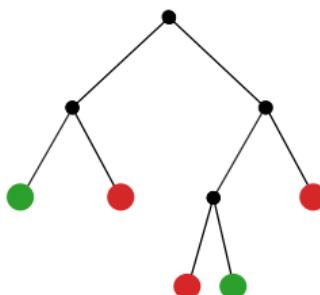
Structure Expansion and controlled Reduction



SER for class imbalance

SER_R, SER_{LL}

Structure Expansion and controlled Reduction

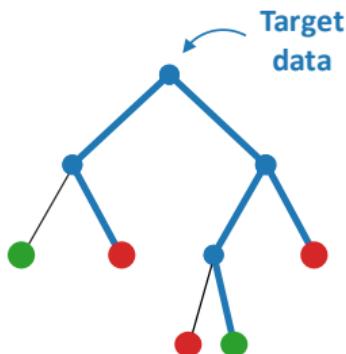


Minority class

SER for class imbalance

SER_R, SER_{LL}

Structure Expansion and controlled Reduction

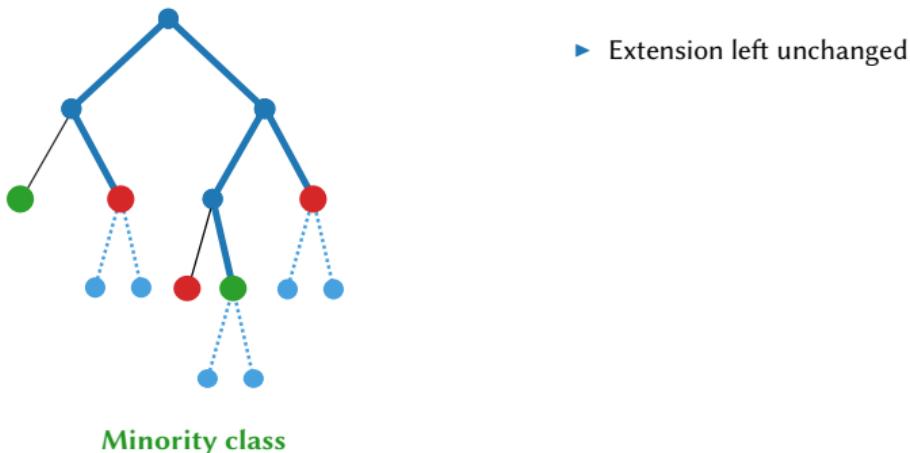


Minority class

SER for class imbalance

SER_R, SER_{LL}

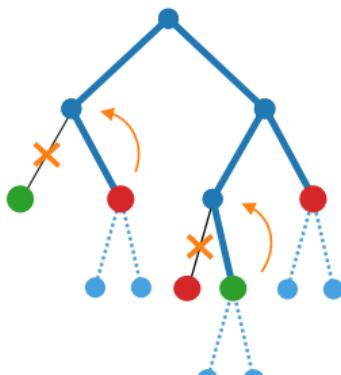
Structure Expansion and controlled Reduction



SER for class imbalance

SER_R, SER_{LL}

Structure Expansion and controlled Reduction



- ▶ Extension left unchanged
- ▶ Reduction constrained

SER_R

If node is of minority class, then no pruning

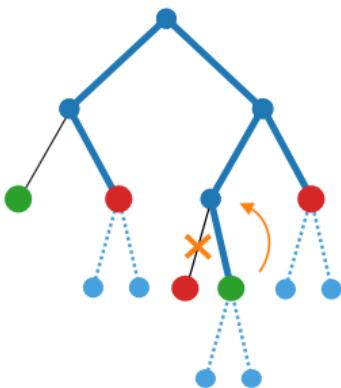
SER_{LL}

If node is of minority class **and** still significant considering Target **and** $R_L > 0.5$, then no pruning

SER for class imbalance

SER_R, SER_{LL}

Structure Expansion and controlled Reduction



Minority class

- ▶ Extension left unchanged
- ▶ Reduction constrained

SER_R

If node is of minority class, then no pruning

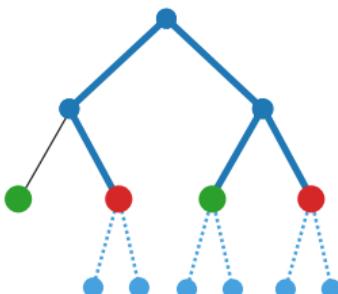
SER_{LL}

If node is of minority class **and** still significant considering Target **and** $R_L > 0.5$, then no pruning

SER for class imbalance

SER_R, SER_{LL}

Structure Expansion and controlled Reduction



Minority class

- ▶ Extension left unchanged
- ▶ Reduction constrained

SER_R

If node is of minority class, then no pruning

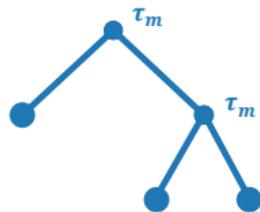
SER_{LL}

If node is of minority class **and** still significant considering Target **and** $R_L > 0.5$, then no pruning

STRUT for class imbalance

STRUT optimization

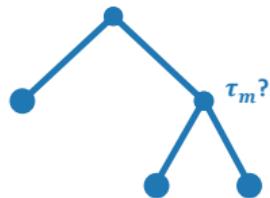
STRUT: How are updated the new thresholds ?



STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?

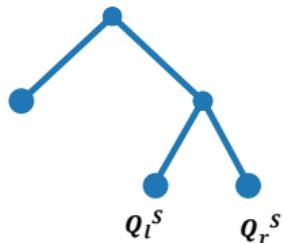


STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?

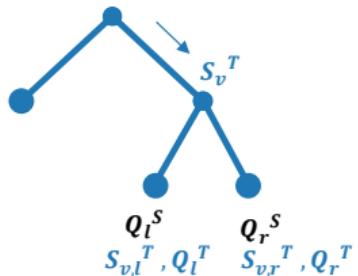
Q_l^S, Q_r^S : class proportions of source data in children w.r.t. the original split



STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?



Q_l^S, Q_r^S : class proportions of source data in children w.r.t. the *original* split

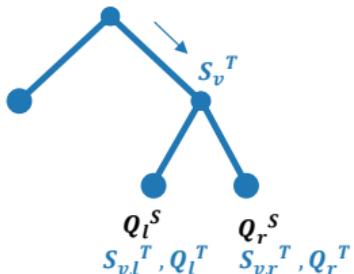
$S_{v,l}^T, S_{v,r}^T$: subsets of S_v^T that fall in the children nodes of v

$Q_l^T(\tau), Q_r^T(\tau)$: class proportions of target data in children w.r.t. the *new* split

STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?



Q_l^S, Q_r^S : class proportions of source data in children w.r.t. the *original* split

$S_{v,l}^T, S_{v,r}^T$: subsets of S_v^T that fall in the children nodes of v

$Q_l^T(\tau), Q_r^T(\tau)$: class proportions of target data in children w.r.t. the *new* split

Divergence Gain: similarity between the original label distributions and the new ones

$$DG(\tau) = 1 - \frac{|S_{v,l}^T|}{|S_v^T|} JSD(Q_l^S, Q_l^T) - \frac{|S_{v,r}^T|}{|S_v^T|} JSD(Q_r^S, Q_r^T)$$

Jensen-Shannon divergence:

$$JSD(P, Q) = \frac{1}{2} (D_{KL}(P||M) + D_{KL}(Q||M))$$

$$M = \frac{1}{2} (P + Q)$$

Kullback-Leibler divergence:

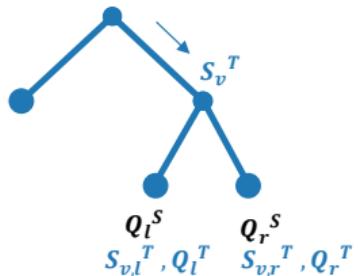
$$D_{KL}(P||Q) = \sum_k P(k) \ln \left(\frac{P(k)}{Q(k)} \right)$$

STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?

Goal: Maximize DG while being in a local maximum of Information Gain (IG) (here IG = Gini gain)

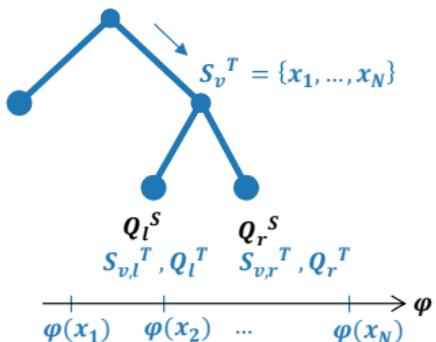


STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?

Goal: Maximize DG while being in a local maximum of Information Gain (IG) (here IG = Gini gain)

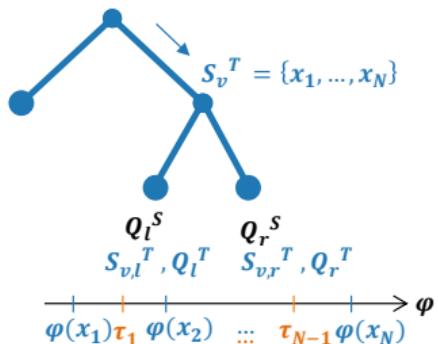


STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?

Goal: Maximize DG while being in a local maximum of Information Gain (IG) (here IG = Gini gain)

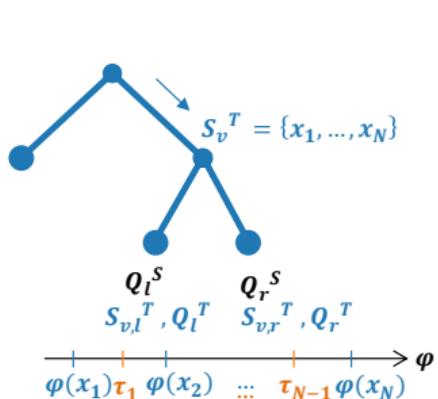


STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?

Goal: Maximize DG while being in a local maximum of Information Gain (IG) (here IG = Gini gain)



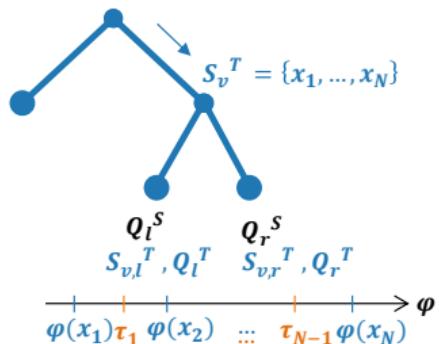
$$\tau_m = \arg \max_{\tau \in T_v} (DG(\tau, Q_l^T(\tau), Q_r^T(\tau)))$$

$$\text{s.t. } IG(\tau_{m-1}) < IG(\tau_m) \text{ and } IG(\tau_m) > IG(\tau_{m+1})$$

STRUT for class imbalance

STRUT optimization

STRUT: How are updated the new thresholds ?



Goal: Maximize DG while being in a local maximum of Information Gain (IG) (here IG = Gini gain)

$$\tau_m = \arg \max_{\tau \in T_v} (DG(\tau, Q_l^T(\tau), Q_r^T(\tau)))$$

$$\text{s.t. } IG(\tau_{m-1}) < IG(\tau_m) \text{ and } IG(\tau_m) > IG(\tau_{m+1})$$

- ▶ Q^S have less meaning when going deeper
- ▶ Do we really want to keep Q^S and Q^T close ?

STRUT for class imbalance

STRUT_{IG}, STRUT_{HI}

STRUT_{IG}

STRUT without DG: maximization of IG

STRUT for class imbalance

STRUT_{IG}, STRUT_{HI}

STRUT_{IG}

STRUT without DG: maximization of IG

STRUT_{HI}

Framework of homogeneous class imbalance, use

$$p^T(y/x) = \lambda_y \frac{p^S(y/x)}{\int \lambda_y p^S(y/x) dy} \quad (2)$$

to change the source class proportions in DG:

$$Q_l^{S'} = \lambda_k \frac{Q_l^S}{\sum_k \lambda_k Q_l^S} \quad Q_r^{S'} = \lambda_k \frac{Q_r^S}{\sum_k \lambda_k Q_r^S}$$

STRUT for class imbalance

STRUT_{IG}, STRUT_{HI}

STRUT_{IG}

STRUT without DG: maximization of IG

STRUT_{HI}

Framework of homogeneous class imbalance, use

$$p^T(y/x) = \lambda_y \frac{p^S(y/x)}{\int \lambda_y p^S(y/x) dy} \quad (2)$$

to change the source class proportions in DG:

$$Q_l^{S'} = \lambda_k \frac{Q_l^S}{\sum_k \lambda_k Q_l^S} \quad Q_r^{S'} = \lambda_k \frac{Q_r^S}{\sum_k \lambda_k Q_r^S}$$

STRUT_{HI} can be seen as a generalization of STRUT

Data

Gaussian Generator

Each dataset consists of a combination of several Gaussian cluster. For each class k with N_k clusters of this class :

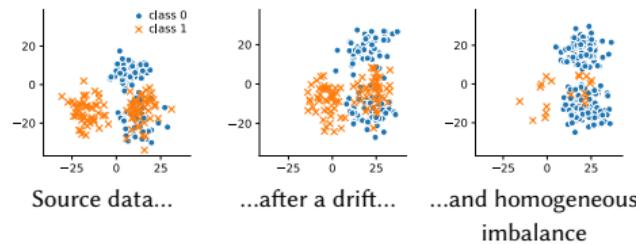
$$C^k = \{\mathcal{N}(\mu_1, \sigma_1), \dots, \mathcal{N}(\mu_{N_k}, \sigma_{N_k})\}$$

$$\forall k, (X, Y = k) \sim \sum_{j=1}^{N_k} w_j^k \mathcal{N}(\mu_j, \sigma_j)$$

Dataset synthetic transformations

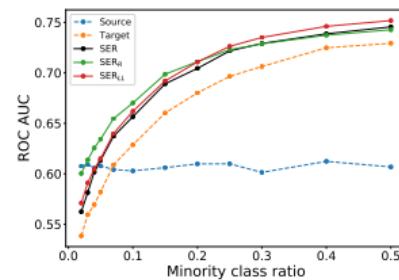
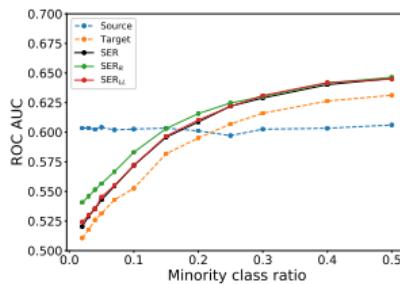
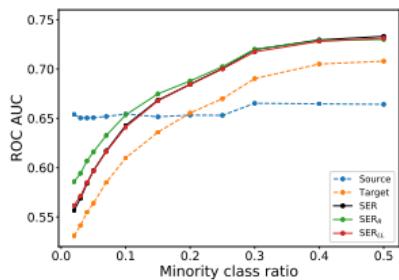
Some basic synthetic transformations on Source clusters are applied to get a Target dataset

- ▶ Imbalance : Clusters weights changes (w_j^k)
- ▶ Drifts (μ_j)
- ▶ Squeezes and Stretches (σ_j)
- ▶ Adding or removing clusters



Results

Synthetic data



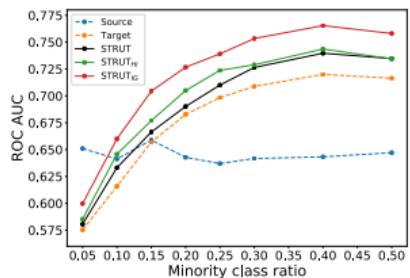
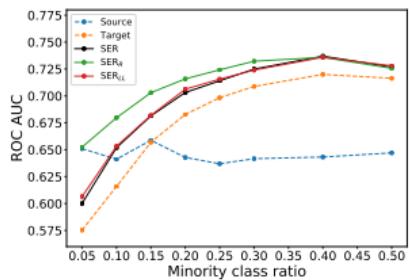
(a) Drift

(b) Stretch / Squeeze

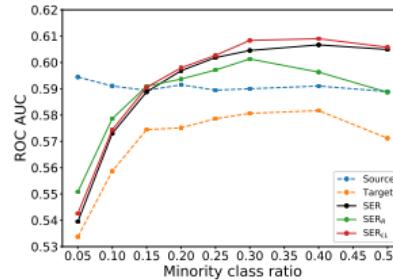
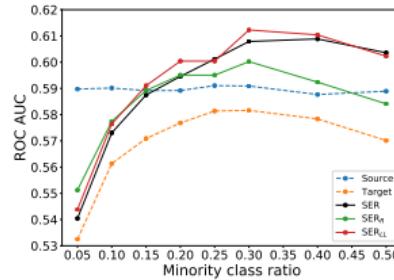
(c) Add / Remove

Results

Real data



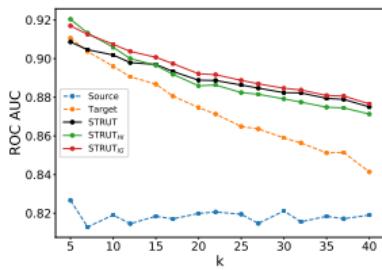
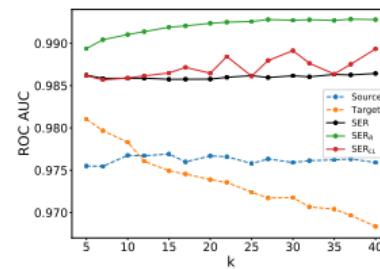
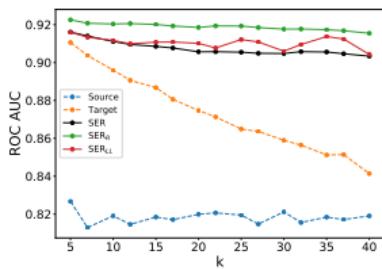
(a) Magic Gamma Telescope

(b) Office-Caltech: *amazon* → *webcam*(c) Office-Caltech: *caltech* → *webcam*

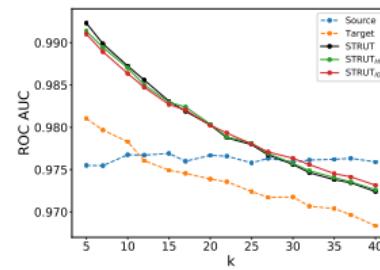
Results

Fall data

Operational data set: 174 fall events and 2619 non-fall events (6%)
k-fold testing with varying k



(a) Decision tree model



(b) Random forest model with 10 decision trees

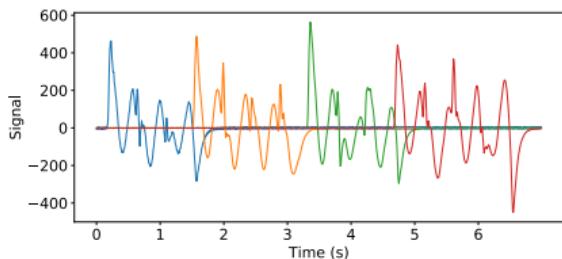
Elderly activity recognition with convolutional representations

Motivation

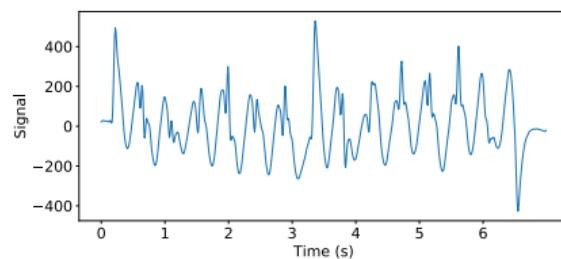
- ▶ **Issue:** one-dimensional signals for large areas
- ▶ Goal: Classify elderly from other individuals
 - ▶ Most signals are made of walks of staff individuals
- ▶ **Subtask:** Bring the model's attention over step-related signals

Motivation

- ▶ **Issue:** one-dimensional signals for large areas
- ▶ Goal: Classify elderly from other individuals
 - ▶ Most signals are made of walks of staff individuals
- ▶ **Subtask:** Bring the model's attention over step-related signals



(a) Raw signal



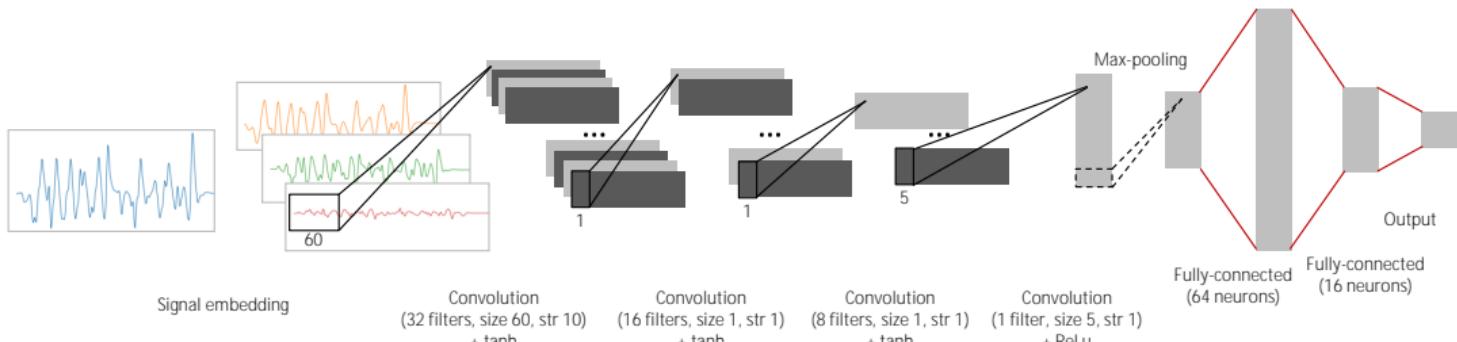
(b) Preprocessed signal

- ▶ Signals are rather complex

This work: A model to recognize elderly activity using convolutional representations and three training steps:

1. Step proposal network inspired from Region proposal network
2. Signal embedding using convolutional dictionary learning
3. The final classification task

Main architecture



- ▶ Use the convolutional representation to “boost” training
- ▶ First layer (Signal embedding) of NURSENET is trained **separately** using convolutional dictionary learning

Signal embedding

Convolutional dictionary learning

- ▶ \mathbf{s} : data to be represented
- ▶ Objective : find M atoms \mathbf{d}_m and activation signals \mathbf{x}_m such that

$$\mathbf{s} \approx \sum_{m=1}^M \mathbf{x}_m * \mathbf{d}_m$$

- ▶ $*$: convolution

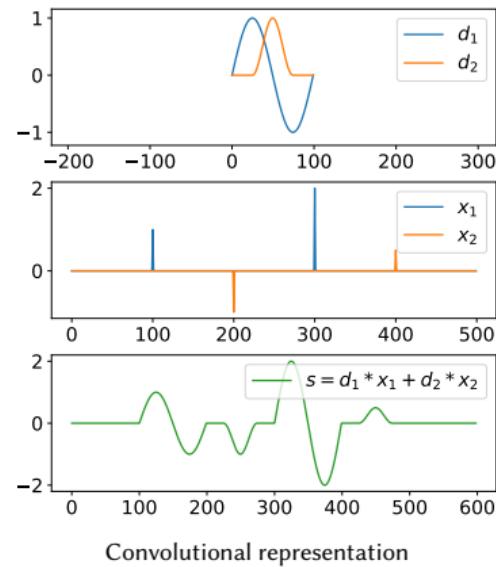
Signal embedding

Convolutional dictionary learning

- ▶ s : data to be represented
- ▶ Objective : find M atoms \mathbf{d}_m and activation signals \mathbf{x}_m such that

$$s \approx \sum_{m=1}^M \mathbf{x}_m * \mathbf{d}_m$$

- ▶ $*$: convolution



Convolutional representation

Signal embedding

Convolutional dictionary learning

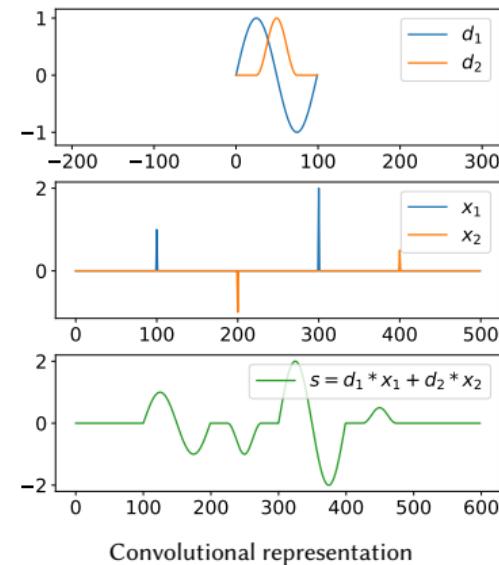
- ▶ \mathbf{s} : data to be represented
- ▶ Objective : find M atoms \mathbf{d}_m and activation signals \mathbf{x}_m such that

$$\mathbf{s} \approx \sum_{m=1}^M \mathbf{x}_m * \mathbf{d}_m$$

- ▶ $*$: convolution

CDL general problem:

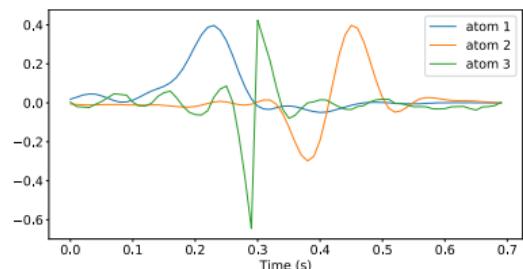
$$\begin{aligned} \arg \min_{\mathbf{x}_m, \mathbf{d}_m} & \frac{1}{2} \left\| \sum_{m=1}^M \mathbf{x}_m * \mathbf{d}_m - \mathbf{s} \right\|_2^2 + \lambda \sum_{m=1}^M \|\mathbf{x}_m\|_1 \\ \text{s.t. } & \|\mathbf{d}_m\|_2 \leq 1 \quad \forall m . \end{aligned}$$



Signal embedding

Learning step atoms

- ▶ Learning with Alternating Direction Method of Multipliers (ADMM) [3]
- ▶ 3 atoms of length 0.7 second

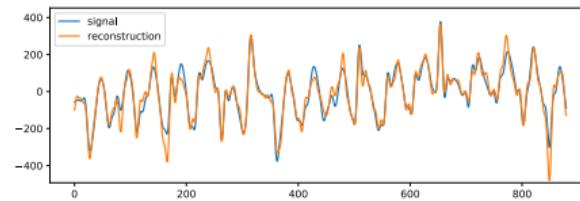
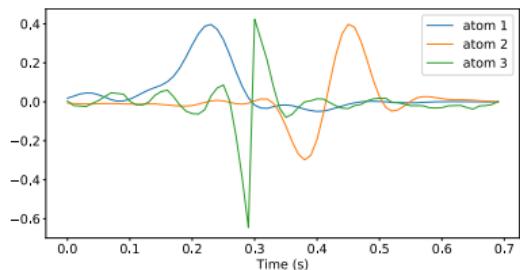


Dictionary

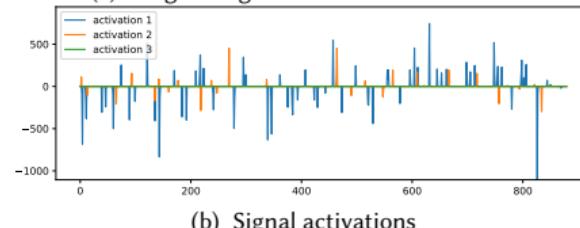
Signal embedding

Learning step atoms

- ▶ Learning with Alternating Direction Method of Multipliers (ADMM) [3]
- ▶ 3 atoms of length 0.7 second



(a) Original signal and its reconstruction



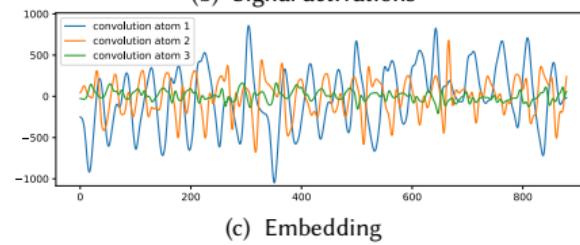
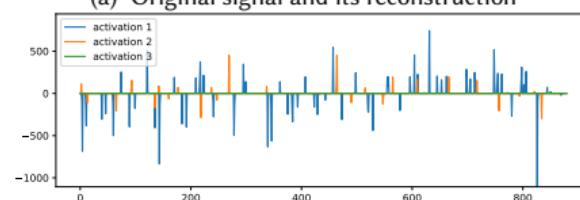
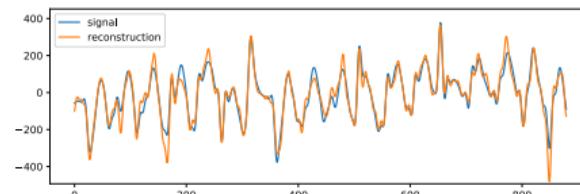
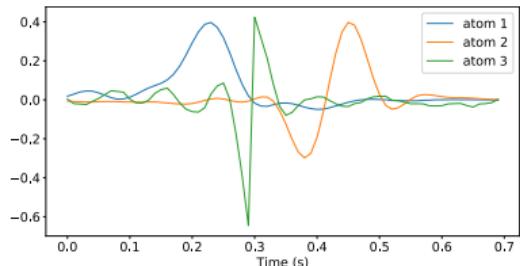
(b) Signal activations

Signal embedding

Learning step atoms

- ▶ Learning with Alternating Direction Method of Multipliers (ADMM) [3]
- ▶ 3 atoms of length 0.7 second
- ▶ Use the following embedding:

$$\mathbf{S} \doteq (\mathbf{s} * \mathbf{d}_m)_{1 \leq m \leq 3}$$



Region proposal network

Object detection

- ▶ Classification: What is the image class ?



Classification vs Object detection. Source: Girshick et al. [4], Ren et al. [6]

Region proposal network

Object detection

- ▶ Classification: What is the image class ?
- ▶ Object detection: Where are the objects and what are they classes ?

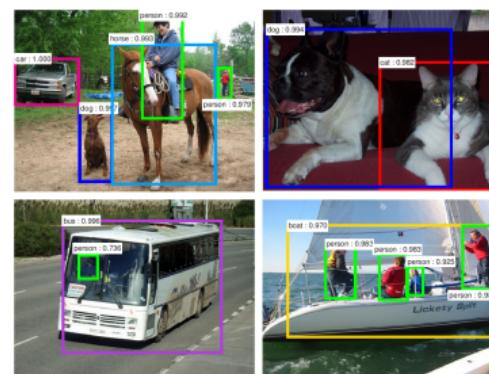


Classification vs Object detection. Source: Girshick et al. [4], Ren et al. [6]

Region proposal network

Object detection

- ▶ Classification: What is the image class ?
- ▶ Object detection: Where are the objects and what are they classes ?
- ▶ How to efficiently localize objects ?
- ▶ Proposal models (Hosang et al. [5])
- ▶ Faster R-CNN (Ren et al. [6])

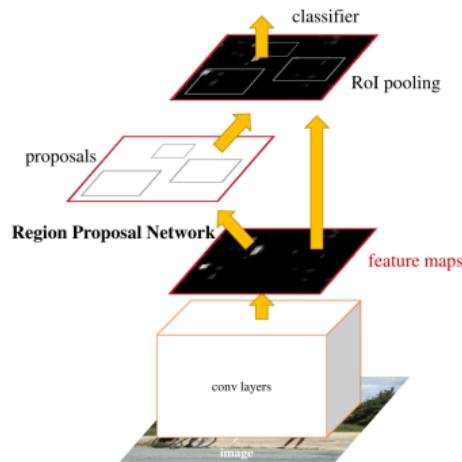


Classification vs Object detection. Source: Girshick et al. [4], Ren et al. [6]

Region proposal network

Faster R-CNN

- ▶ Main idea: proposals are generated by a CNN called Region Proposal Network

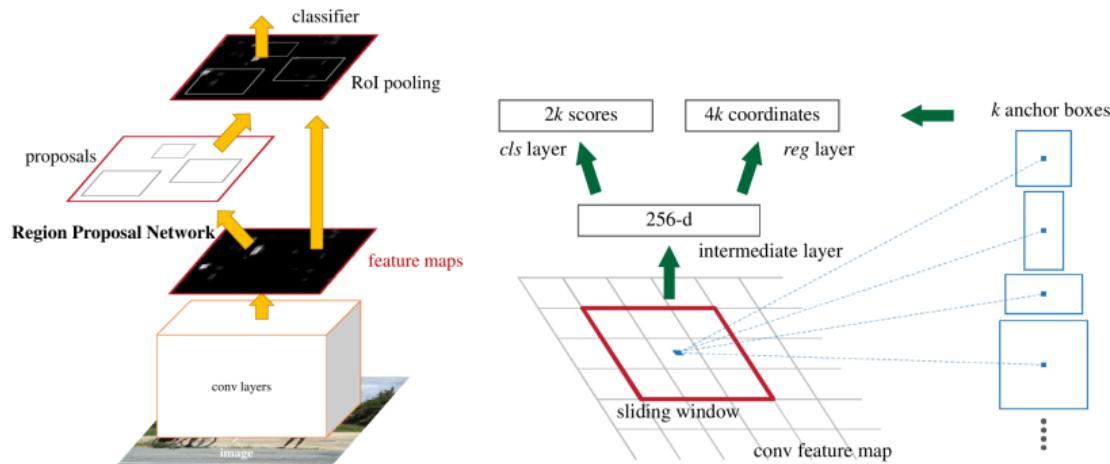


Region proposal network. Source: Ren et al. [6]

Region proposal network

Faster R-CNN

- ▶ Main idea: proposals are generated by a CNN called Region Proposal Network
- ▶ A sliding window is passed: multiple *anchors* over each location (various sizes and scales)
- ▶ Two layers: Classification (Object / Not Object) and Regression (anchor coordinates)

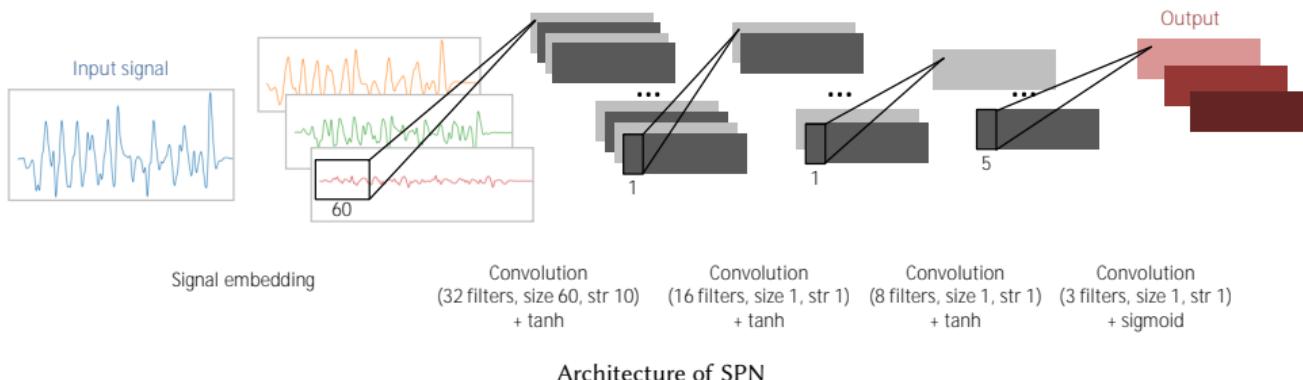


Region proposal network. Source: Ren et al. [6]

Step proposal network

Main architecture

- ▶ Directly inspired from RPN
- ▶ Simple architecture with three hidden layers, all **convolutional**
- ▶ Output: probability of having a step at a specific window location and size
 - ▶ Here 3 sizes and all discrete locations are considered

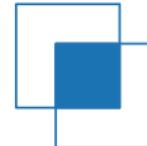


Step proposal network

Principle

- ▶ Objective of SPN : output boxes with largest Intersection over Union (IoU)
- ▶ IoU: \mathbf{b}_j are labelled boxes, \hat{b} is an estimated box:

$$\text{IoU}(\hat{b}) \doteq \max_j \frac{|\mathbf{b}_j \cap \hat{b}|}{|\mathbf{b}_j \cup \hat{b}|}$$

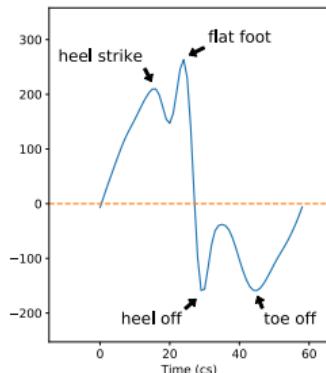


Step proposal network

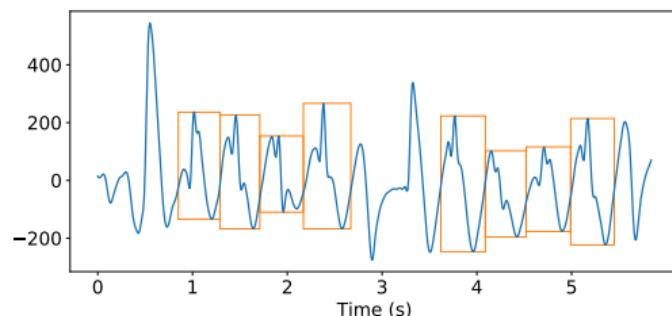
Principle

- ▶ Objective of SPN : output boxes with largest Intersection over Union (IoU)
- ▶ IoU: b_j are labelled boxes, \hat{b} is an estimated box:

$$\text{IoU}(\hat{b}) \doteq \max_j \frac{|\mathbf{b}_j \cap \hat{b}|}{|\mathbf{b}_j \cup \hat{b}|}$$



(a) Step signal



(b) Step labels over a walk signal

Step proposal network

Training

- ▶ Output: a matrix $\mathbf{W} \in \mathbb{R}^{T \times K}$
 - ▶ T : signal length
 - ▶ K : number of different box sizes
- ▶ $\mathbf{W}_{t,k}$: probability that the box b_t^k starting at time t and of size 0.4s, 0.5s, or 0.6s (for respectively $k = 1, 2$, or 3) has a large IoU score
- ▶ Positive boxes: $\text{IoU}(b_t^k) > \sqrt{0.7}$
- ▶ Negative boxes: $\text{IoU}(b_t^k) < \sqrt{0.3}$
- ▶ Other are not used for training

The loss function \mathcal{L} over a signal \mathbf{s} is defined as:

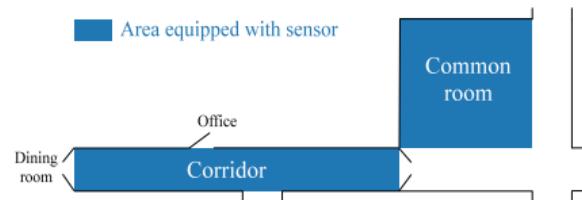
$$\mathcal{L}(\mathbf{s}, \mathbf{W}) = \sum_t \sum_{k \in [1, 2, 3]} \mathbb{1}_{\text{IoU}(b_t^k) > \sqrt{0.7}} \log(\mathbf{W}_{t,k}) + \mathbb{1}_{\text{IoU}(b_t^k) < \sqrt{0.3}} \log(1 - \mathbf{W}_{t,k}).$$

Step proposal network

Results

Data

- ▶ 43 signals recorded in a nursing home
- ▶ Manually labeled steps

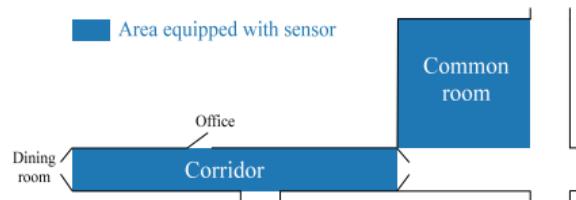


Step proposal network

Results

Data

- ▶ 43 signals recorded in a nursing home
- ▶ Manually labeled steps



Training

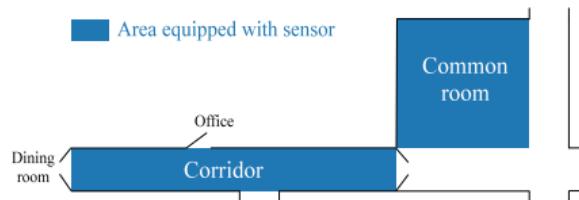
- ▶ SPN is trained using classical gradient descent
- ▶ Training time: < 5 minutes
- ▶ Inference (detection over a 10s signal): < 1 second
- ▶ Optimization details
 - ▶ learning rate of 10^{-3}
 - ▶ learning rate decay ($\times 0.9$ every 10 epochs)
 - ▶ Nesterov momentum

Step proposal network

Results

Data

- ▶ 43 signals recorded in a nursing home
- ▶ Manually labeled steps



Results

- ▶ Object detection use the mean Average Precision (mAP): area under the Precision-Recall curve
- ▶ **Without** embedding, mAP = 72,5%
- ▶ **With** embedding, mAP = 78,6%

Training

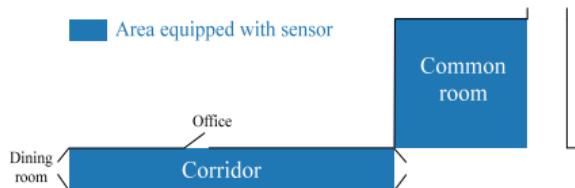
- ▶ SPN is trained using classical gradient descent
- ▶ Training time: < 5 minutes
- ▶ Inference (detection over a 10s signal): < 1 second
- ▶ Optimization details
 - ▶ learning rate of 10^{-3}
 - ▶ learning rate decay ($\times 0.9$ every 10 epochs)
 - ▶ Nesterov momentum

Step proposal network

Results

Data

- ▶ 43 signals recorded in a nursing home
- ▶ Manually labeled steps

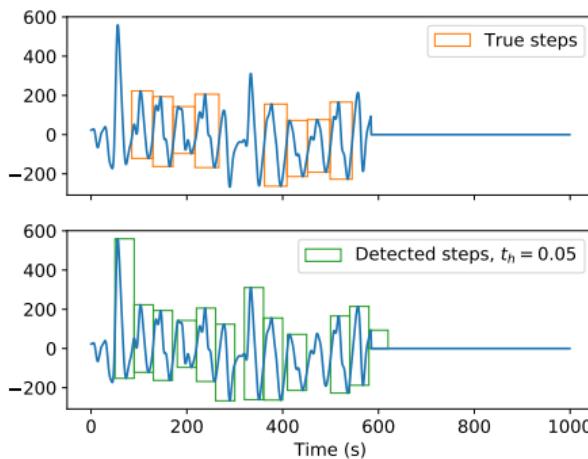


Training

- ▶ SPN is trained using classical gradient descent
- ▶ Training time: < 5 minutes
- ▶ Inference (detection over a 10s signal): < 1 second
- ▶ Optimization details
 - ▶ learning rate of 10^{-3}
 - ▶ learning rate decay ($\times 0.9$ every 10 epochs)
 - ▶ Nesterov momentum

Results

- ▶ Object detection use the mean Average Precision (mAP): area under the Precision-Recall curve
- ▶ **Without** embedding, mAP = 72,5%
- ▶ **With** embedding, mAP = 78,6%

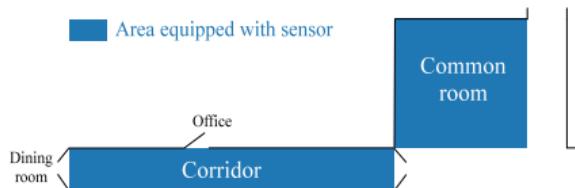


Step proposal network

Results

Data

- ▶ 43 signals recorded in a nursing home
- ▶ Manually labeled steps

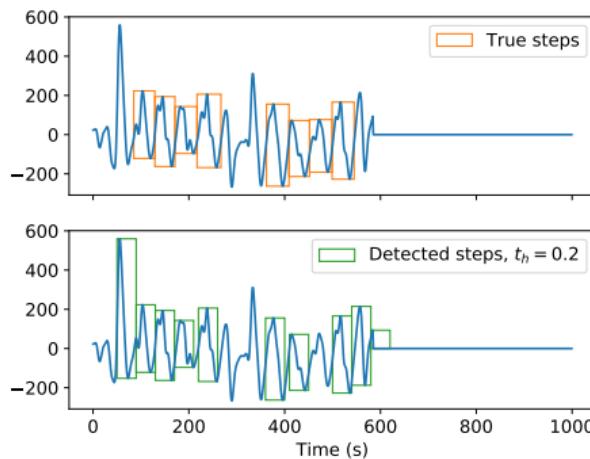


Training

- ▶ SPN is trained using classical gradient descent
- ▶ Training time: < 5 minutes
- ▶ Inference (detection over a 10s signal): < 1 second
- ▶ Optimization details
 - ▶ learning rate of 10^{-3}
 - ▶ learning rate decay ($\times 0.9$ every 10 epochs)
 - ▶ Nesterov momentum

Results

- ▶ Object detection use the mean Average Precision (mAP): area under the Precision-Recall curve
- ▶ **Without** embedding, mAP = 72,5%
- ▶ **With** embedding, mAP = 78,6%

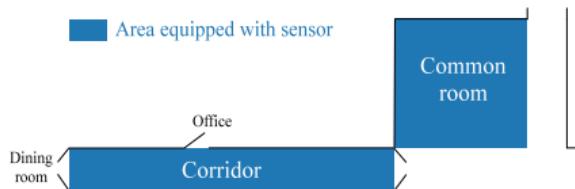


Step proposal network

Results

Data

- ▶ 43 signals recorded in a nursing home
- ▶ Manually labeled steps

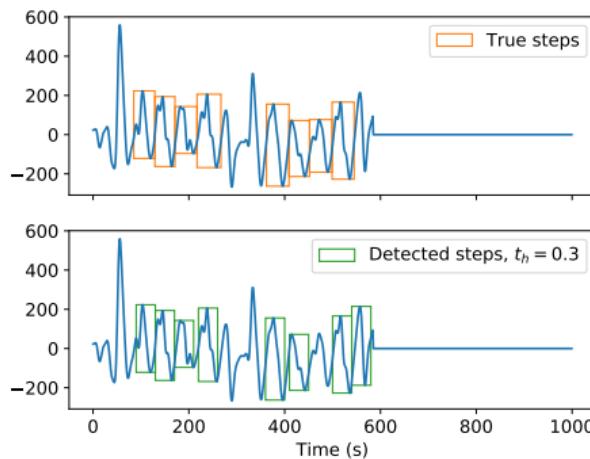


Training

- ▶ SPN is trained using classical gradient descent
- ▶ Training time: < 5 minutes
- ▶ Inference (detection over a 10s signal): < 1 second
- ▶ Optimization details
 - ▶ learning rate of 10^{-3}
 - ▶ learning rate decay ($\times 0.9$ every 10 epochs)
 - ▶ Nesterov momentum

Results

- ▶ Object detection use the mean Average Precision (mAP): area under the Precision-Recall curve
- ▶ **Without** embedding, mAP = 72,5%
- ▶ **With** embedding, mAP = 78,6%



Classification

Training

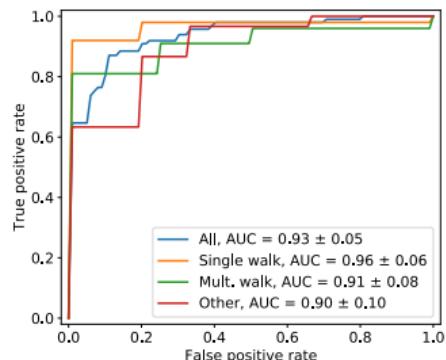
$$\mathcal{L}(\mathbf{s}_i, \hat{y}_i) = -(\mathbb{1}_{y_i=1} \log(\hat{y}_i) + \mathbb{1}_{y_i=0} \log(1 - \hat{y}_i))$$

Sublabel	Staff	Elderly
Single walk	43	31
Multiple walks	19	11
Other	19	23

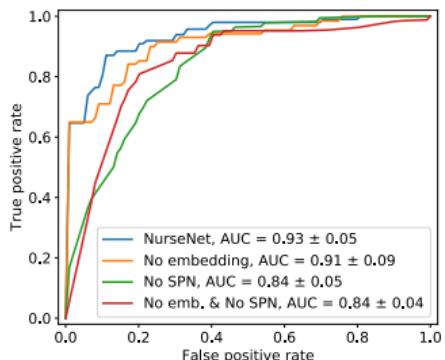
- ▶ SPN is trained using classical gradient descent
- ▶ Training time: < 5 minutes
- ▶ Inference (detection over a 10s signal): < 1 second
- ▶ Optimization details
 - ▶ learning rate of 10^{-3}
 - ▶ learning rate decay ($\times 0.9$ every 10 epochs)
 - ▶ Nesterov momentum

Classification

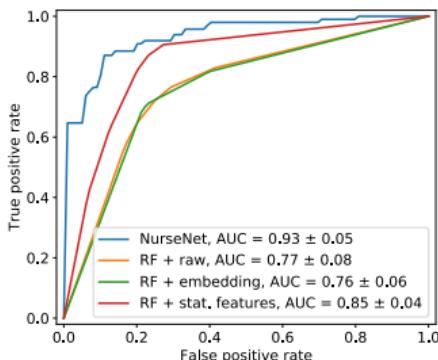
Results



(a) NURSENET results



(b) NURSENET ablation analysis



(c) NURSENET and several RF-based models

Conclusion

Contributions

- ▶ A simple and practical model for fall detection
- ▶ Transfer procedure for decision tree adapted to class imbalance
- ▶ A model to distinguish elderly vs. other with high accuracy

Future work

- ▶ Explore meta models for transfer procedures
- ▶ What if features change ? Consider heterogeneous transfer
- ▶ Improve step detection precision (with regression layer ?)
- ▶ Test on benchmarks step data sets
- ▶ Is it possible to distinguish activities or individuals ?

Publications and communications

Publications

- ▶ L. Minvielle, M. Atiq, R. Serra, M. Mougeot, and N. Vayatis. [Fall detection using smart floor sensor and supervised learning](#). In *2017 39th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, pages 3445–3448, July 2017
- ▶ L. Minvielle, M. Atiq, S. Peignier, and M. Mougeot. [Transfer learning on decision tree with class imbalance](#). In *2019 IEEE 31st International Conference on Tools with Artificial Intelligence (ICTAI)*, pages 1003–1010, Nov 2019
- ▶ Ludovic Minvielle and Julien Audiffren. [Nursesnet: Monitoring elderly levels of activity with a piezoelectric floor](#). *Sensors*, 19(18), 2019
- ▶ P. Humbert, B. Le Bars, L. Minvielle, and N. Vayatis. [Robust Kernel Density Estimation with Median-of-Means principle](#). *Submitted to Neural Information Processing Systems 2020 (NeurIPS)*, 2020

Communications

- ▶ Transfer learning on decision tree with imbalanced data, 3rd Summer school on transfer learning, 2019, Passau, Germany
- ▶ Step detection with proposal network. 4th French-German Summer School on Artificial Intelligence, 2020, Online.

References

- [1] Leo Breiman. Random forests. *Machine Learning*, 45(1):5–32, 2001.
- [2] Leo Breiman, Jerome H. Friedman, Richard A. Olshen, and Charles J. Stone. *Classification and regression trees*. Monterey, CA: Wadsworth & Brooks/Cole Advanced Books & Software, 1984.
- [3] Hilton Bristow, Anders Eriksson, and Simon Lucey. Fast convolutional sparse coding. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2013.
- [4] R. Girshick, J. Donahue, T. Darrell, and J. Malik. Rich feature hierarchies for accurate object detection and semantic segmentation. In *2014 IEEE Conference on Computer Vision and Pattern Recognition*, pages 580–587, 2014.
- [5] J. Hosang, R. Benenson, P. Dollár, and B. Schiele. What makes for effective detection proposals? *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 38(4):814–830, 2016.
- [6] Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun. Faster r-cnn: Towards real-time object detection with region proposal networks. In *Advances in neural information processing systems*, pages 91–99, 2015.
- [7] N. Segev, M. Harel, S. Mannor, K. Crammer, and R. El-Yaniv. Learn on source, refine on target: A model transfer learning framework with random forests. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 39(9):1811–1824, Sep. 2017.